

Dissertation

Tyler Wiederich

2025-09-25

Table of contents

2	CHAPTER 0	5
2.1	Motivation and Background	5
2.2	Overview of Physical 3D Visualizations	8
2.2.1	Creating 3D-Printed Visualizations	9
2.3	Bar Charts in Research and Practice	11
2.3.1	Against 3D Bar Charts	12
2.3.2	For 3D Bar Charts	13
2.3.3	Conflicting Empirical Studies	15
2.4	Heat Maps in Research and Practice	19
2.4.1	Empirical Studies of 2D and 3D Heat Maps	21
2.5	Methods for Evaluating Visualization Effectiveness	23
2.5.1	What makes a good chart?	23
2.5.2	Designing studies for numerical estimations	24
3	CHAPTER 1	27
3.1	Introduction	27
3.1.1	Elementary Graphical Tasks	28
3.1.2	3D Graphical Perception	30
3.2	Methods	33
3.2.1	Replicating Cleveland and McGill	33
3.2.2	Stimuli Construction	34
3.2.3	Experiment Design	37
3.2.4	Participant Recruitment	37
3.2.5	Data Collection	39
3.3	Results	41
3.3.1	Midmeans of Log Absolute Errors	42
3.3.2	Generalized Additive Model	43
3.3.3	Interactivity with 3D Rendered Charts	47
3.4	Discussion and Future Work	48
3.4.1	Supplemental Material	51
4	CHAPTER 2	52
4.1	Abstract	52
4.2	Introduction	53
4.3	Methods	55
4.3.1	Overview and Motivation	55
4.3.2	Experiential Learning	55
4.3.3	Data Analysis	57
4.3.4	Graphics Experiment	58

4.4	Results	60
4.4.1	Recruitment of Students and Instructors	60
4.4.2	Selected Responses from Experiment Participation	60
4.4.3	Selected Responses from Research Reflections	62
4.5	Discussion	63
4.6	Figure legends	66
4.7	Tables	67
5	CHAPTER 3	70
5.1	Introduction	70
5.1.1	Literature Review	71
5.2	Methods	73
5.2.1	Stimuli	73
5.2.2	Charts	76
5.2.3	Subject Recruitment	79
5.2.4	Experimental Design	80
5.3	Limitations	85
5.4	Conclusion	86
6	REFERENCES	87

1

This is the non-chapter stuff

2 CHAPTER 0

2.1 Motivation and Background

Data visualizations play an essential role in understanding patterns in the structure of data. They allow for a programmatic mapping of data into a “picture” that can be used to gather insight from the viewer (Tukey 1965; Tufte 2001; Wilkinson and Wills 2005). Since the early 20th century, researchers have been exploring the question: what makes a good graph? (Croxtton and Stryker 1927a; Croxtton and Stein 1932a; Cleveland and McGill 1984a; Vanderplas, Cook, and Hofmann 2020b). Many attempts to answer this question have provided good recommendations, but are largely limited to the projection of graphs onto 2-dimensional (2D) surfaces. While this addresses many of the typical use cases for data visualizations, the process of creating modern statistical graphics in our 3-dimensional (3D) world is a relatively new phenomenon that does not have widespread usage.

Many charts exist within the confines of a 2D space due to the practicality of computer renderings. These computer-generated graphics are quick and cheap to produce (Tukey 1965), which possibly contributes to their widespread usage. Before computers, hand-drawn and complex printing techniques were common production methods and took time to produce (Friendly and Wainer 2021; VanderPlas, Ryan, and Hofmann 2019). In contrast, all data

physicalizations require additional material resources, increasing the base cost to produce. Additionally, there are few modern resources that have direct pipelines from data to physical objects. One such example is the rayshader package (Morgan-Wall 2024), although modifications to the resulting stereolithography (STL) file or OBJ file are needed before physicalization can take place. This results in the majority of modern graphics being produced with 2D representations on digital screens.

Nearly every type of chart can be created by using a programmatic mapping of data and variables to aesthetics and geometries, a process sometimes referred to as the “Grammar of Graphics” (Wilkinson and Wills 2005). This has found widespread implementation into numerous software (Wickham 2010; Satyanarayan et al. 2017; “SAS 9.4 ODS Graphics: Procedures Guide, Sixth Edition,” n.d.). For example, consider the chart created in Figure Figure 2.1. In this figure, the Palmer Penguins dataset (Horst, Hill, and Gorman 2020) is supplied to the `ggplot` function from the `ggplot2` package (Wickham 2016). The x and y axes are mapped to bill length and bill depth, with color and shape being mapped to island. The geometry is given by `geom_point`, which results in a scatter plot that shows each complete data observation. Although the chart would benefit from further customization, it was quick to produce and has instant visual insights. Other software packages can create the same chart using a similar process by specifying the axes, colors, and shapes using a point geometry.

```
library(tidyverse)

library(palmerpenguins)

ggplot(penguins, mapping = aes(x = bill_length_mm,
                               y = bill_depth_mm,
```

```

    color = island,
    shape = island)) +
geom_point()

```




Figure 2.1: Example of a graphic created using the grammar of graphics.

While there have been many attempts to establish good data visualization practices (Tuft 2001; Vanderplas, Cook, and Hofmann 2020b), we do not know how well these recommendations translate from 2D projections into our 3D world. The advent of 3D printing allows us to precisely construct the 3D charts that have widely been studied through their 2D representations (Croxtan and Stein 1932a; Barfield and Robless 1989a; Zacks et al. 1998). With limited research on 3D-printed data manifestations, we explore the initial empirical findings of true 3D statistical graphics.

2.2 Overview of Physical 3D Visualizations

Data can be expressed in the physical world through a wide variety of means. In many cases, this can take on artistic expressions, such as encoding information into crocheted blankets or wooden sculptures (Huron et al. 2023). While certainly providing information about the data, these informational art pieces are largely left as display items to promote conversation into the underlying dataset.



Figure 2.2: 3D-printed treemap created displaying Gross National Income (GNI) for 130 countries (Huron et al. 2023, pg. 217). The two 3D prints on the bottom have additional information overlaid on top of the treemap as 2D displays.

Perhaps one of the most affordable types of data physicalization is through the use of 3D printing. Aside from the startup cost of a 3D-printer, material and labor costs are relatively low. Rolls of PLA filament can be acquired for under \$20 and can be used to create multiple charts. Print times vary depending on the size and complexity of the graphic, but smaller charts can be produced in approximately 2 hours. While needing much more time to


produce than computer-generated visualizations, the amount of hands-on labor during the 3D-printing process is minimal.

One of the earliest 3D-printed charts was shared in 2014 by the Wall Street Journal, where the print was used to discuss the rising popularity of 3D printers (Graphics n.d.; Thingiverse.com n.d.). Since then, using 3D printers to create statistical graphics has typically been reserved for niche purposes. It is rare to find these types of charts in practice, where images of the charts are often posted to social media or blog posts. This results in limited access to the physical charts, even though source print files can be shared through sites such as [thingiverse.com](https://www.thingiverse.com) and [printables.com](https://www.printables.com).


2.2.1 Creating 3D-Printed Visualizations

The popularity of the 3D printer exploded in the early 2010's as the technology became increasingly available and cheaper. This has led to increased use of 3D printing in many focus areas, including healthcare (Dodziuk 2016) and engineering (V et al. 2023), but has only seen novel use cases for statistical graphics. Many software programs have the ability to create 3D statistical graphics, but lack the ability to easily export the graphics into files suitable for 3D printing.

Numerous software programs can create digital renderings of 3D charts. Excel (Microsoft Corporation 2018) has native support for creating charts with 3D depth cues, but requires add-ins to produce charts using 3 axes. R (XXX) has several options for creating 3D charts that have data on 3 axes (e.g., Murdoch and Adler 2023; Morgan-Wall 2024; Sievert 2020), but lacks support for creating depth charts. SAS 9.4 (XXX) has options such as PROC GCHART for adding depth cues and PROC G3D for surface charts. Other popular



(a) Original 3D printed chart from WSJ



(b) Replicated print created in 2025

Figure 2.3: 3D-printed volume charts showing the sales of 3D printers under \$5,000 (Thingiverse.com n.d.). Figure 2.3a shows the original WSJ print from 2014. Figure 2.3b shows two prints created with a Flash Forge Finder 3D Printer, where the smaller chart took 2 hours to print and the larger chart took 6 hours to print.

software programs contain similar capabilities, such as JavaScript libraries, MATLAB, and Python.

While a number of tools exist for creating 3D data visualizations, the pipeline of getting these charts into files compatible with 3D printing is not widely automated. For example, R has some options using the `rgl` (Murdoch and Adler 2023) and `rayshader` (Morgan-Wall 2024) packages, but both packages have limited tuning parameters for the resulting output files. Another option is to use 3D software, such as OpenSCAD (Kintel 2023), to manually create the charts, but these software programs lack the ability to directly integrate statistical information in the formulation of the charts.

Another consideration for 3D-printed charts is the inclusion of text labels. With single-filament 3D printers, text can be incorporated in three primary ways: (1) embossing, in which the text is raised above the surface of the print; (2) engraving, in which the text is recessed into the surface; or (3) applying external labels, such as stickers or ink. Multi-filament 3D printers have the additional option to use a different color to make text labels flush with their surfaces. However, Munzner (n.d.) cautioned that the text of 3D charts will suffer from legibility issues due to the distortion of the text when viewed at an angle.

2.3 Bar Charts in Research and Practice

One of the most common types of charts is the bar chart, characterized by the use of rectangular geometries in a 2D space. Two variables are mapped to the geometries of the bar chart, where one axis is dedicated to a categorical variable of nominal or ordinal nature, and the other axis is for a continuous variable to denote magnitude or a response. These

charts have been used over 200 years and remain a popular choice for many modern data visualizations (Friendly and Wainer 2021).

The use of 3-dimensional elements in bar charts typically involve converting the rectangular bars into rectangular prisms. Figure 2.4 shows an illustration of this transformation.




Figure 2.4: Construction of depth cues for a 2D bar chart. First, the amount of depth is set by defining the opposite face of the standard 2D rectangular shape. Then the faces are connected and volume is filled in.

The naming convention for these types of bar charts is not consistent. Many researchers in the visualization community call these 3D bar charts due to the portrayal of three-dimensional space (e.g., Zacks et al. 1998; Fischer 2000). Other times, the lack of information in the third dimension results in the charts referred to as 2.5D charts (e.g., Tractinsky and Meyer 1999). To keep a consistent nomenclature, they will be called “3D bar charts” throughout this dissertation.

2.3.1 Against 3D Bar Charts

Tufte (2001) called the use of 3D elements a “fake perspective” in his description of “chart junk”, which is where visual elements add clutter to the data visualization. Other researchers (Zacks et al. 1998; Stewart, Cipolla, and Best 2009) have termed the perspective

effect “extraneous” when referring to depth cues. The reasoning behind this is fairly intuitive: why include additional noise when simpler alternatives exist?

The seminal work of Cleveland and McGill (1984a) theorized that encoding information in volumes would provide worse numerical estimation than for encodings using positions, lengths, and areas. Their reasoning stems from Steven’s Power Law (XXX), a psychophysics formulation that estimates a perceived stimuli magnitude p with the actual magnitude a by $p = ka^\alpha$. The estimation of α provides guidance to what types of stimuli are subject to the most distortion to their true scale when $\alpha = 1$. While the 3D bar charts do not lose their 2D encodings, they do gain depth, and thus volume encodings. The effect of volume versus area comparisons was noted earlier by Croxton and Stein (1932a), where accuracy for ratio estimations was worse for cubes as compared to bars and circles.

From an artistic point of view, depth cues can obscure some elements of the chart, resulting in a decrease in readability. For 3D bar charts, this often occurs from the volumetric elements partially covering the response axis. Consider Figure 2.5 where the y-axis corresponds to the back face of the rectangular prism. In this case, the front and right faces do not directly correspond with the axis, thus making it more difficult to extract values from the chart.

2.3.2 For 3D Bar Charts

Levy et al. (1996) argued that depth cues can be useful in certain contexts. In three experiments involving over 100 psychology students, participants were asked to select from a range of 2D and 3D charts for different tasks. Results showed that students more often preferred 3D charts when the goal was to enhance memory of the chart, while preferences




Figure 2.5: Figure in Statistical Graphics: Design Principles and Practices. The viewing angle for this chart would have viewers overestimate the values of the bars

were about evenly split between 2D and 3D when the task was to present information to others. Levy et al. (1996) suggested that the depth embellishments in 3D graphics may boost memorability by making charts more visually distinctive.

This idea aligns with broader findings on “chart junk,” where non-data elements have been shown to aid recall of bar charts, though often at the cost of longer processing times (Bateman et al. 2010; Borgo et al. 2012; Peña, Ragan, and Harrison 2020). Similar increases in processing time have been reported specifically for 3D bar charts (Siegrist 1996; Fischer 2000; Stewart, Cipolla, and Best 2009), lending some support to the claim that depth cues may improve memorability despite potential cognitive trade-offs.

The persistence of 3D bar charts reflects the influence of artistic preferences in visualization. The added depth can make results appear more striking, giving the impression of visual weight. Such embellishments often appeal to audiences by enhancing visual impact, even when they add little to the clarity of the data itself.

2.3.3 Conflicting Empirical Studies

Although Tufte (2001) argued against the use of 3D charts when possible, empirical studies involving direct comparisons of 2D and 3D bar charts have shown mixed results. In general, these studies focus on metrics of accuracy and find that 3D bar charts tend to be either less accurate or just as accurate as their 2D counterparts. Another common metric is response time, where 3D charts tend to have longer response times. Overall preference for one style of chart over another is a subjective measurement, but has shown mixed results for the two styles of charts.

2.3.3.1 Accuracy

Accuracy is a common metric in studies comparing 2D and 3D bar charts, typically measured by either extracting a single numeric estimate or making comparisons between two values. Many studies used controlled environments for constructing stimuli, often displaying one or two stimuli at a time, or controlling non-target stimuli. In practice, charts are often more complex, displaying many stimuli without drawing attention to particular values.

When only measuring accuracy, 3D bar charts tend to be less accurate than their 2D counterparts (Zacks et al. 1998; Stewart, Cipolla, and Best 2009). However, this result is less prominent when introducing time delays (Zacks et al. 1998, Experiment 2) or easier task complexity (Stewart, Cipolla, and Best 2009). Siegrist (1996) noted that bar positions and sizes also affected performance accuracy, emphasizing that other factors contribute to accuracy. It is also important to note that metrics of accuracy are not conclusive, where sometimes there is no evidence of a difference between 2D and 3D bar charts (Melody Carswell, Frankenberger, and Bernhard 1991).

Of course, there are other less-studied factors that could affect perceptual judgments of accuracy in 3D bar charts. In addition to viewing parameters found in 2D bar chart studies (Fischer, Dewulf, and Hill 2005; Rice et al. 2024), adding depth cues are also subject to viewing angles, amount of depth, and field of view parameters. Viewing angle, with respect to the axis, can mislead the reader into overestimating or underestimating numerical quantities Figure 2.5. Zacks et al. (1998) did not find a significant effect for the amount of perception depth, but noted a trend such that increased exaggeration of depth lowered their accuracy metric. Lastly, field of view is fixed for the viewer, but can widely distort 3D renderings Figure 2.6. These factors add additional complexity, but are not widely studied

for the use of statistical graphics.




Figure 2.6: Multiple charts of the volcano dataset (R Core Team 2024) rendered using the `rgl` package (Murdoch and Adler 2023) in R. Each chart has a different field-of-view parameter, with the amount of distortion increasing as the field-of-view (FOV) increases.

2.3.3.2 Response Times

The speed in making perceptual judgments is another common metric in studies comparing types of charts. In theory, charts where questions can be answered more quickly implies that the chart is better at communicating information. However, response time is a poor metric in situations of exploratory data analysis and long term interactions with complex graphics (Vanderplas, Cook, and Hofmann 2020b). When a question of interest is known

in advance, response time becomes a more viable metric.

Depth cues from 3D bar charts tend to increase the amount of time required to answer questions (Siegrist 1996; Fischer 2000; Stewart, Cipolla, and Best 2009). Intuitively, additional complexity should require additional time to process. For 3D bar charts, viewers need to align the response axis with the rectangular prisms in order to determine the approximate magnitude of the bar. Consider Figure X, where the y-axis denotes magnitude with respect to a hidden face of the bar. In order to extract a numerical quantity, the viewer first has to align the axis with the correct face of the prism before processing the magnitude.

2.3.3.3 Preference


The debate between whether to use 2D or 3D graphics stems from a wider adoption of 3D charts in practice. Computer graphics have made it easy to produce charts of either dimensionality, resulting in creative liberties for the creation of the chart for publication. Figure 2.7 showcases two examples of 3D bar charts from the 1978 Handbook of Agricultural Charts (*1978 Handbook of Agricultural Charts. Enlargements / [u.s. Department of Agriculture] 1978*) where other options exist for displaying the data. Schmid (1983) (Ch. 8) also showcases many other 3D charts published in official reports.

Preference can be broken down into two broad categories: (1) visual appeal, and (2) usefulness. Visual appeal is where perceptual tasks are not a priority, and the goal is to have an image that is pleasant to the viewer. Usefulness, on the other hand, is a rating of how effective the viewer thought the chart was in order to answer the question. Both are subjective measures that influence the choice

When not tasked with estimating values, 3D bar charts tend to be more visually appealing

CHANGES IN CANNED VEGETABLE CONSUMPTION PER CAPITA BETWEEN 1970-72 AND 1975-77


TOTAL LB.



Fresh-weight basis.

CHANGES IN CANNED VEGETABLE CONSUMPTION PER CAPITA BETWEEN 1970-72 AND 1975-77

TOTAL LB.



Fresh-weight basis.

(a) Bar charts with strictly positive values


(b) Bar chart with negative values

Figure 2.7: Two examples of 3D bar charts found in publication (*1978 Handbook of Agricultural Charts. Enlargements / [u.s. Department of Agriculture] 1978*). This left image is a stacked bar chart, using shadings to indicate sugarbeets and sugarcane. The image on the right

(Levy et al. 1996; Fisher, Dempsey, and Marousky 1997a; Stewart, Cipolla, and Best 2009). This may be in part due to the forced illusion of depth which creates additional embellishments that would increase the impression and memorability of the chart, as seen by other types of Tufte’s “chart junk” (Borgo et al. 2012; Peña, Ragan, and Harrison 2020). Despite the visual appeal, 2D charts tend to be chosen for preference and ease of use when making numerical estimations (Levy et al. 1996; Stewart, Cipolla, and Best 2009).

2.4 Heat Maps in Research and Practice

When creating charts for three or more variables, it becomes increasingly complex to visualize data (Grinstein and Trutschl 2001). One option is to map the additional variables using gestalt principles (Vanderplas, Cook, and Hofmann 2020b) (Original citation?), which



(a) 3D volume chart
Three examples of chart preferences.

includes making use of color, size, and/or shapes to denote information. However, not every aesthetic is compatible with every type of variable. Color can be used to map either continuous variables with gradient scales or categorical variables with palettes. Size is typically reserved for continuous variables, where larger values correspond to a larger quantity. Shapes tend to have finite distinguishable representations in order to represent categorical variables. Examples of good and poor cases for each of these aesthetics is presented in Figure Figure 2.11.

Heat maps are a type of chart that can be used with two variables of any type and a continuous response variable. In two-dimensional formats, color is often used to denote the magnitude of the response, utilizing a gradient scale. When represented in three dimensions, height is the primary mapping of the response of the heat map. Due to the height aesthetic, 3D heat maps are sometimes referred to as “height maps”, particularly when spatial coordi-




Figure 2.11: Examples of color, size, and shape for scatterplots of the penguins dataset (Horst, Hill, and Gorman 2020).


nates are involved. Since the color height of the response variable can only be mapped once to each combination of x and y coordinates, heat maps are generally used on the summary of responses, such as averages or single observations.

2.4.1 Empirical Studies of 2D and 3D Heat Maps

Unlike 2D and 3D bar charts, the comparison of dimensionality for heat maps has not been widely studied. This may in part be due to the complex nature of color scales (Breslow, Ratwani, and Trafton 2009; Liu and Heer 2018b; Reda and Papka 2019). While bar charts simply add a perspective depth cue to increase the dimensionality, the translation of color to height results in numerous considerations, such as the scaling of the height Figure 2.13.


In 2D heat maps, the choice of color palettes can influence the perception of numerical

Participation in Postsecondary Education of the High School Class of 1972: Fall 1974




(a) 2D heat map

Participation in Postsecondary Education of the High School Class of 1972: Fall 1974



(b) 3D heat map

Figure 2.12: Heat maps created from with data from (XXX). The 2D heat map represents the percentage of high school graduates using a gradient color scale, whereas the 3D heat map uses height.



(a) Flatter height scale



(b) Stretched height scale

Figure 2.13: Two figures representing the volcano dataset (R Core Team 2024) using different height scaling factors. Both figures use the same color scale, but there is no “correct” translation of color into the measurable height.

estimations (Liu and Heer 2018a; Molina Lopez, Middel Soria, and Vázquez 2023; Zhang et al. 2023). Some color palettes are easier to extract information from than other. An additional consideration is the use of colorblind-friendly palettes for people with color deficiencies. This relationship between color and 2D heat maps contributes to the difficulties with direct translations to 3D heat maps.

FIGURE PLACEHOLDER: COLOR PALETTES WITH COLORBLIND FILTERS

Only a few studies have directly compared the dimensionality of statistical graphics (Kraus et al. 2020; Jeong et al. 2025). Conducted primarily in virtual reality, these studies suggest that 2D and 3D charts each offer advantages depending on the task. Two-dimensional charts generally produced higher interpretation accuracy and better recognition of overall distributional patterns. In contrast, three-dimensional charts were more effective for estimating single values and supporting long-term recall. Taken together, these findings highlight dimensionality as an important consideration when selecting visual displays for specific analytic goals.

2.5 Methods for Evaluating Visualization Effectiveness

2.5.1 What makes a good chart?

Vanderplas, Cook, and Hofmann (2020b) discussed various approaches to testing graphics, including numerical estimations, speed, and error rates. These metrics are fairly common metrics when comparing 2D and 3D charts (Croxtton and Stein 1932a; Cleveland and McGill 1984a; Hughes, n.d.), but do not represent the full picture of what makes a one type of chart better than another. As seen by Levy et al. (1996), the objectives of the chart can determine

which type of chart should be used. While many metrics exist, the practical limitation of testing generally allows for only a few metrics to be tested at a given time.

Many studies involving statistical graphics tend to rely on numerical estimations for recommending data visualization practices. This is particularly useful given the role that data visualizations play in the understanding of underlying datasets (Tukey 1965).

2.5.2 Designing studies for numerical estimations

Over 2,500 citations have referenced the work of Cleveland and McGill (1984a) when discussing properties of statistical graphics. Since then, the methods of analyzing and conducting studies on statistical graphics have evolved (e.g., Vanderplas, Cook, and Hofmann 2020b; Hofmann et al. 2012), but the structure of experimental designs involving numerical estimations tends to remain similar. Tasks typically involve the comparison of two stimuli, where the question is “how much bigger/smaller is one value compared to the other?” Care needs to be taken in the phrasing of this question since mental subtractive processes are perceived differently from mental ratio processes (Veit, n.d.; Hagerty and Birnbaum 1978).

Cleveland and McGill (1984a) asked participants two questions when making comparisons between two values on a chart: (1) which of the two values was the smaller, and (2) what percentage the smaller was of the larger. The first question clarified if the participants were looking at the values correctly for the second question, and the second question is an estimation of A/B , where $A < B$. Accuracy was measured by $\log_2(|\text{True value} - \text{Estimated value}| + 1/8)$, using the trimmed averages when bootstrapping confidence intervals. Cleveland and McGill’s experiment was replicated by Heer and Bostock (2010a), showing similar results

for the experiment’s procedure in a crowdsourced environment.

When designing studies for numerical estimation, the choice of A/B and the measure of error are important considerations. Many researchers seemingly choose ratios in an arbitrary fashion or follow previously established work Melody Carswell, Frankenberger, and Bernhard (1991). Examples of ratios used in several studies are provided in Figure 2.14.

As for measures of error, several formulations can be used:

- $|\text{Error}|$ for magnitude of error
- Error for direction of error
- $\log |\text{Error}|$ for non-linear error estimations

Psychophysics has shown that different estimation tasks affect how participants respond to numerical estimations. Estimations of ratios tend to fall on an exponential scale, where larger differences in stimuli result in smaller differences of estimated ratios. This makes $\log |\text{Error}|$ a natural choice for measuring ratio estimation tasks, similar to how Cleveland and McGill (1984a) measured error. Measurements of differences tend to be estimated linearly, where estimated differences are the same regardless of differences in stimuli. As long as the estimation task is consistent for each trial, statistical methodology can handle the differences in estimation tasks.




Figure 2.14: Ratios used by several studies

3 CHAPTER 1

3.1 Introduction

Good communication requires both that the information be transmitted correctly and that the intended recipient be able to decode and understand the transmitted information accurately. In order to communicate effectively, we must use graphical forms that accurately convey information relevant to the task in question. In many cases, this means we must understand how accurately people can read quantitative information off of charts. While accuracy is not the only quantity of interest in graphical investigations (Hullman et al. 2019), it is an important factor in assessing the utility of many different data graphics.

The accuracy of graphical forms has been studied for almost a century (von Huhn 1927; Eells 1926a; Croxton and Stein 1932b; Croxton and Stryker 1927b); as new ways of representing information evolve, we must revisit old studies to determine whether these representations have the same limitations as previous versions. This is particularly true in areas like graphics which are affected by the immense technological innovation in hardware and software which has taken place since the early 1990s.

3.1.1 Elementary Graphical Tasks

Cleveland and McGill (1984b) established the comparative accuracy of different “elementary perceptual tasks” (EPTs). Elementary Perceptual Tasks, according to these experiments, include assessing graphical elements such as position along a common scale, length, angle, and volume, and estimating the corresponding numerical value of these representations. The study relied entirely on estimation accuracy, which may not always be relevant when extracting information from graphs. For example, estimation is less relevant when ordering values by size. As a result of the Cleveland and McGill (1984b) study, it is possible to assemble an ordering of perceptual accuracy for the elements of length, position, and angle. Heer and Bostock (2010b) replicated some parts of Cleveland and McGill (1984b) in an online setting using Mechanical Turk, a platform for crowdsourcing human tasks that has been used for (among other things) assembling machine learning datasets. This replication study largely validated the results of the original study while demonstrating the utility of the Mechanical Turk platform for graphical testing.

The first experiment in Cleveland and McGill (1984b) (the position-length experiment), used five types of bar charts: two types of grouped bar charts and three types of stacked bar charts. Each chart had two bars marked for comparison; participants were asked to determine which bar was smaller and give the perceived ratio of the smaller bar to the larger bar. **Figure 1** shows the two types of grouped bar charts. We are primarily interested in the grouped bar charts (in part because 3D printing is not yet inexpensive enough to make moderate-scale stacked bar chart experiments viable), which consisted of two comparison bars which were either adjacent or in separate groups. These grouped bar charts will be referenced as adjacent and separated graph types in this paper, respectively.





Figure 3.1: Cleveland and McGill (1984b) used two different types of grouped bar charts: comparisons between adjacent bars, and comparisons between separated bars. It is widely acknowledged that comparisons between separated bars (Type 3 comparisons, in Cleveland & McGill's terminology) are more difficult and error-prone.

Figure 3.2: Cleveland and McGill (1984b) used two different types of grouped bar charts: comparisons between adjacent bars, and comparisons between separated bars. It is widely acknowledged that comparisons between separated bars (Type 3 comparisons, in Cleveland & McGill's terminology) are more difficult and error-prone.


3.1.2 3D Graphical Perception

Chart perception is affected by the human visual system’s implicit assumption that visual stimuli are three-dimensional; after all, most of the visual input we process does come from a three-dimensional world, but charts are artificial. In addition, these artificial stimuli are typically created and presented in two dimensions. The artificial scenes we create for data visualization have been documented to cause problems: for instance, the line-width illusion (VanderPlas and Hofmann 2015a; Day and Stecher 1991; Hofmann and Vendettuoli 2013) has been attributed to implicit 3D perception of 2D stimuli and can affect perception of error bands, candlestick plots, hammock plots, and Sankey diagrams. The line width illusion causes lines or error bands which are actually the same length to appear shorter at points with large amounts of curvature; a simple example of the classic sine illusion (one version of the line width illusion) is shown in Figure 3.3. While this illusion is problematic in two dimensions, when the depth cues are actually present in a three-dimensional situation, the same implicit visual heuristics contribute to accurate size perception (as in the middle panel of Figure 3.3). Thus, there is reason to think that perceptual accuracy may be dependent on the realism of the visual display relative to the training of the visual system and the heuristics which are activated in order to make sense of the display.


The use of 3D graphics have been explored in multiple studies, with mixed results. Fisher, Dempsey, and Marousky (1997b) explored user preference for 2D or 3D charts and found that subjects tended to prefer simpler 2D graphs when tasked with extracting information. Barfield and Robless (1989b) compared 2D and 3D graphs presented on paper and on computers, showing that the accuracy of subject answers depended on their skill level: novice subjects were more accurate with 2D paper graphs, while experienced managers were more



(a) Sine Illusion



(b) 3D surface with similar features



(c) Perceived vs. actual length

Figure 3.3: An illustration of the sine illusion (VanderPlas and Hofmann 2015a), also known as the line-width illusion. All vertical lines are the same length, but the lines in the middle of the curve appear to be much shorter. The illusion results when implicit perceptual corrections useful for perceiving the size of objects with depth are applied to 2D stimuli with no actual depth. As 3D heuristics occasionally cause inaccurate communication when applied to 2D objects, it is reasonable to think that there might be some situations where charts making use of a realistic third dimension might be perceived more accurately than their 2D equivalents.

accurate with 3D computer graphs. For both experience levels, participants were more confident in their answers when using 2D graphs. There are instances where 2D graphs perform better than 3D graphs, but there are times where 3D graphs may provide a more natural way to encode information, at least in theory. For instance, when X and Y are used to represent spatial dimensions, it may be preferable to use a 3D chart to convey numerical information instead of using color, which is perceived much less accurately. Brath (2014) highlights the intrinsic attributes of 3D graphs and the benefits when used appropriately with other 3D elements such as lighting and correct portrayal of data attributes. Nevertheless, it is extremely common in visualization literature to recommend that 3D plots be avoided at all costs (Vanderplas, Cook, and Hofmann 2020a; Wilke 2020; Morgan-Wall 2019).

There is thus good reason to be wary of the use of three dimensions where only two are necessary to convey data (Huff 1954). However, the situation is different now than it was in 1984 when Cleveland & McGill published their seminal work; it has even changed since Heer’s replication study (2010b) in 2010.

Computer graphics and rendering technology have developed quickly, along with the hardware necessary to support these software developments. As a result, we have much more natural virtual rendering of 3D objects, and we can also print graphics in three dimensions using mass-market, relatively inexpensive 3D printers, moving artificial charts into a more natural, physical setting. As a result, it is reasonable to reconsider the use of 3D charts, not only because of new technological developments, but also because 3D printed charts could eventually provide the opportunity to make data graphics accessible to those with limited or absent vision (Fleet 2023).

In this paper, we discuss a study designed to examine Cleveland & McGill’s experiments

on grouped bar charts using modern graphics in two and three dimensions. This study lays the groundwork for additional empirical studies on the use of 3D graphics, rendered and 3D printed, for visualizing complex data. In order to explore perception of fully 3D graphics, it is prudent to start with the simplest possible 3D graphic: one in which the third dimension is not necessary, so that we can easily compare to two-dimensional representations without loss of information. In the next section we provide details about design and execution of our experiment, including the process of replicating stimuli from Cleveland and McGill (1984b) in order to create 2D, 3D projected, 3D rendered, and 3D printed bar graphs. The next section presents the results, and we conclude the paper by discussing this experiment in the context of existing work on the perception of 2D and 3D graphical elements.

3.2 Methods

Our study is designed to replicate and expand upon the position-length experiment from Cleveland & McGill as closely as possible, with the hope of being able to integrate this study with previously reported results in a consistent manner. In this section, we discuss the replication process and the design of our version of this experiment.

3.2.1 Replicating Cleveland and McGill

The first step of replicating the position-length experiment was to determine the heights of the bars that participants use for comparisons. These values for the bar heights are linear on a log scale and are given by

$$[s_i = 10^{(i-1)/12}], \quad i=1, \dots, 10]$$

Each graph presents two bars from the values given above where the participants are asked to judge the ratio of the smaller bar to the larger bar. The ratio of bars used by Cleveland and McGill were 17.8, 26.1, 38.3, 46.4 (twice), 56.2, 68.1 (twice), and 82.5 (twice). The exact numeric comparisons were not disclosed, but the comparison values used in our study were subjected to the constraints of having the same ratio values and that no value was used more than twice.

Table 3.1: Values used in the experiment to make ratio comparisons, sorted by ratio. The ID label corresponds to the file reference number.

Bar	ID: 01	ID: 06	ID: 05	ID: 04	ID: 02	ID: 03	ID: 09
Smaller	10.0	14.7	14.7	12.1	10.0	12.1	26.1
Larger	56.2	56.2	38.3	26.1	17.8	17.8	31.6
Ratio (%)	17.8	26.2	38.4	46.4	56.2	68.0	82.6

In each graph, there are two sets of five bars each, and two of the bars are marked for identification with a circle and triangle, as in [?@fig-bar-types](#). Cleveland and McGill did not specify the random process for the heights of the eight other bars, so we used a scaled Beta distribution to provide random structure around our fixed values. Code to reproduce the data generation process, data underlying the plots used in this study, the rendered plots and STL files, anonymized user data, and analysis code can found be at <https://github.com/TWiedRW/2023-JDS-3dcharts>.

3.2.2 Stimuli Construction

The graphs share a common layout across all formats, where two groupings of five bars are identified by “A” and “B”, respectively, and circles and triangles are used to identify

the bars participants should compare. Example graphs are shown in `?@fig-plotTypes`. There are some graphical elements that cannot be easily portrayed via 3D printing. For this reason, all graph types do not have axes, grid lines, or floating titles.




Figure 3.4: Two dimensional, 3D digital rendering, and 3D-printed charts used in this study.




Figure 3.5: Two dimensional, 3D digital rendering, and 3D-printed charts used in this study.




Figure 3.6: Two dimensional, 3D digital rendering, and 3D-printed charts used in this study.

The `ggplot2` (Wickham 2016) package was utilized to create the 2D bar charts. The scale axis was removed, leaving only the bars and a bar grouping identifier. The bars used for comparisons had the identifying mark at a height of 5 out of 100 for the 2D plots, and the 3D plots had the identifying marks on top of the bars.

The 3D renderings and 3D printed charts were both created using OpenSCAD (Kintel 2023), which creates STL files from markup describing the object's geometric composition. Charts were composed of a platform, with raised text labels centered in front of the A

and B groups of bars. Bars were created by inserting bar heights into the OpenSCAD template using R, and then raised circle and triangle markers were positioned on top of the bars to be compared to provide a tactile and visual indicator. In addition, an ID code was engraved into the bottom of the platform to uniquely identify each object; this allowed the researchers to ensure that each kit contained the correct charts throughout the experiment while minimizing the risk of stickers or other identifiers becoming detached from the charts. As printed, the base of each chart was 13cm x 3cm x 1cm, with the highest bar rising 9.5cm above the chart base. Raised letters and shapes were 2mm above the base or bar, respectively. The color of the filament used to print each chart corresponded to the estimation ratio (though this was not obvious to participants), allowing the researchers to quickly determine that a kit of 3D charts contained no ratio duplication. Colors were randomly assigned to ratios so that participants could not gain information about the ratio from any perceived ordering of the colors.

Digital renderings of the generated STL files were created using RGL (Murdoch and Adler 2023), which integrates into Shiny (Chang et al. 2024) using the WebGL (Mozilla Foundation 2023) extension. The RGL rendering of the STL file was initially angled to correspond to the perspective of default 3D bar charts present in Microsoft Excel, but in the WebGL interactive environment, participants could rotate the charts as they pleased to create a useful estimate of the bar ratio. Rendered charts were colored to correspond with the filament color used to print the physical chart for consistency. The default `rgl` lighting was replaced with three lights located in fixed positions around the rendered figure. The lights were positioned so that one was behind the rendered figure, another in front of the figure, and one light below the figure. This arrangement provided a consistent visual experience and minimized shadows that prevented visual estimation of the bars.

3.2.3 Experiment Design

A diagram of the experimental design is provided in Figure 3.7. The study is set up as a modified randomized incomplete block design, comprised of 3 blocks (presentation mode) and 7 bar ratios (5 are selected for each participant). At each combination of ratio and presentation type, the comparison type, adjacent (Type 1) or separated (Type 3), is randomly selected. This process was used to create 21 kits of 5 3D printed charts and 10 virtual charts. There is thus a slight deviation from a completely randomized incomplete block design, in that all possible combinations of ratios, presentation modes, and comparison types are not present in our design. However, this approach provided the best experimental design given the constraints of managing physical stimuli; it took nearly a month of continuous 3D printing to create all of the charts for the kits that we assembled. A dataset containing the kit composition breakdown is available in the github repository for this project.

3.2.4 Participant Recruitment

One interesting facet of Cleveland and McGill (1984b) is the participant recruitment methodology: “For each experiment the subjects fell into two categories: (1) a group of females, mostly housewives, without substantial technical experience; (2) a mixture of males and females with substantial technical training and working in technical jobs. Most of the subjects in the position-length experiment participated in the position-angle experiment; in all cases repeat subjects judged the position-angle graphs first.” It seems likely that the authors recruited individuals within their respective departments as well as their wives. In order to replicate this atypical participant recruitment method in the modern era, members




Figure 3.7: A graphical representation of the study design. Only five of the seven ratios were used in each kit; at least one of the smallest or largest ratios were included along with 4 other charts; each selected ratio was displayed in all 3 mediums. For each medium $\times$ ratio combination, the comparison type (separated or adjacent) was randomly determined. A total of 21 kits of 3D printed charts were created to include all combinations of the five ratios.

of the UNL Statistics department and their spouses, partners, and roommates were asked to participate in our study. This replicates the spirit of the recruitment method without the implicit assumptions that graduate students and professors are male, heterosexual, and have unemployed spouses. A total of 38 participants completed our study; demographics are shown in Figure 3.8.

While the results of any sample with recruitment methodology like this are not generalizable to the public, running our initial study on a comparable population to that of the initial study serves a purpose: other replications, such as Heer and Bostock (2010b), used online samples of paid participants, who likely have less experience using charts and graphics and making numerical estimations than those in a statistics department (and individuals who cohabitate with them). We have every intention of running this experiment in other




Figure 3.8: Demographic characteristics of participants in the study.

populations, but the convenience of online sampling is not compatible with physical objects such as our 3D printed charts, so we will have to recruit in-person participants. Thus, an initial study in a population comparable to Cleveland and McGill (1984b) is a reasonable first step into testing 3D charts.

3.2.5 Data Collection

A Shiny applet was used for data collection, along with the provided kit of 3D printed charts. Participants provided informed consent through the applet, and then were asked for demographic information (age, gender, education level). Then, participants were shown a “practice” page which allowed them to experiment with the data collection interface and practice estimating the ratio between the bars, as shown in Figure 3.9. This practice interface did not provide any numeric feedback as to participant correctness, but was merely intended to familiarize the participants with the questions which would be asked as well as the process of estimating the ratio between the two bars.

After the practice screen, participants were asked to provide the kit ID. Participants were also instructed that when indicated, they should choose a 3D chart from the kit and select




Figure 3.9: Screenshot of Shiny application practice screen. Three 2D bar charts with different ratios were provided, along with sliders indicating the correct proportion. Participants could practice with the sliders and preview the questions that would be asked as part of the task.

the ID code of the chart (inscribed on the bottom) from the drop-down menu before completing the estimation task. Participants were also instructed to make quick judgements for each graph and not to measure or estimate ratios using physical objects.

Each graph (or prompt, in the case of 3D printed charts) in the applet had two corresponding questions for participants to answer: first, participants were to identify the smaller bar by shape, and then, participants were to estimate the ratio of the size of the smaller bar to the size of the larger bar, as shown in Figure 3.10. It should be noted that we chose to use a slider without numeric identifiers or tick marks for entering ratio estimates - in part, this was intended to alleviate rounding effects which occur when participants enter numbers on their own or get numeric feedback from slider entries (Ruud, Schunk, and Winter 2014; Maineri, Bison, and Luijckx 2021). Another benefit of this modification is that it should reduce participants' cognitive load, allowing them to focus on the ratio between bars rather than the numerical translation process. This differs from Cleveland and McGill (1984b),

Which bar is smaller?

Triangle (▲)

Approximately what size is the smaller bar in comparison to the larger bar?

Smaller Larger

Next




Figure 3.10: Screenshot of the applet collecting data for a 3D rendered chart task. Participants were asked to select which bar (circle or triangle) was smaller, and then to estimate the ratio of the smaller bar to the larger bar.

but maintains the spirit of the study while using methods which were not available for data entry at the time.

3.3 Results

The participants provided a total of 495 responses. While each participant was provided with 15 trials, 6 participants partially completed the experiment; one participant did not advance beyond the demographic page. All responses that incorrectly identified the smaller bar ($n = 18$) were removed from the study before analysis; this serves as a basic attention check and leaves 477 valid responses. In addition, one response was removed because the participant did not enter a valid ID for a 3D printed chart.

3.3.1 Midmeans of Log Absolute Errors

Cleveland and McGill (1984b) used

$$\log_2(|\text{Judged Percent} - \text{True Percent}| + 1/8)$$

to measure accuracy of their participant's responses. In their study, log base 2 seemed appropriate due to "average relative errors changing by factors less than 10." They also added 1/8 to prevent distortions when the errors were close to zero. Heer and Bostock (2010b) followed the same analysis method, replicating many (but not all) of the results presented in the original paper.

Figure 3.11 shows the midmeans of the log absolute errors compared to the true ratio of the bars for each graph type and comparison type. The results suggest that the log absolute errors increase for greater differences between the smaller and larger bars, but are consistent across the graph types.




Figure 3.11: Midmeans and observed values of log absolute errors for the true ratio of bars. Summary lines are computed from raw data using a loess smooth.

This is somewhat different than the results in Cleveland and McGill (1984b) and Heer and

Bostock (2010b); in both cases the midmean log absolute errors increased until about 55% of the true proportional difference and then decreased. It is possible that this difference is due to the fact that we did not require an explicit numerical estimate of the ratio but instead asked participants to indicate the ratio on a slider (which is essentially a number line). It is possible that this change explains the lack of a difference between chart modalities we see, but there are other potential explanations as well. One interesting facet of this data is that there is a consistent reduction in log error when the true proportional difference is near 50%. This may be because of implicit anchoring - we can typically bisect a span relatively accurately, which means that it might be expected that when the true proportion is near 50%, we would both notice that and be able to accurately transfer that information to the slider. This pattern is more noticable in separated (Type 3) comparisons, which also makes sense - as these comparisons are known to be less accurate, a locally minimal error near 50% with increased error relative to adjacent (Type 1) comparisons could easily explain the exaggerated trends we see in Figure 3.11.

3.3.2 Generalized Additive Model

In addition to replicating the (primarily graphical) analysis of participant errors, we also took a more statistical approach and fitted a generalized additive model (Hastie 2017) that accounts for participant variation as well as the effect of comparison type, graph type, and ratio.

We define the following effects:

- R , the numerical ratio of A/B where A is the length of the smaller bar and B is the length of the larger bar. There are $i = 1, \dots, 7$ levels of R in this study, but we will

treat R as a numeric variable throughout this model.

- D_j , the display mode, where $j = 1, 2, 3$ correspond to 2D, 3D rendered, and 3D printed charts respectively.
- T_k , the type of comparison, where $k = 1, 2$ correspond to adjacent (Type 1) and separated (Type 3) comparisons, respectively
- P_l , a random effect for participant which describes the overall skill of the participant at visual estimation of ratios.

Generalized additive models depend on a smooth function $s() : \mathbb{R} \rightarrow \mathbb{R}$ which is applied to the independent variables; by default, the smooth function is created using thin-plate regression splines. In place of an interaction term when a categorical variable is involved, separate smooths are fit for each level of the categorical variable; we will indicate this by subscripting the smooth function with the categorical variable index, so that $s_1^T(R)$ indicates the smooth over numerical variable R for the 1st level of the categorical variable T, that is, the smooth function over R for adjacent comparisons.

Then the formal statistical model is as follows:

$$y = s_j^D(R) + s_k^T(R) + D_j + T_k + P_l + \epsilon$$

where

- $y = \log_2(|\text{Judged Percent}_{ijkl} - \text{True Percent}_{ij}| + 1/8)$
- $s_j^D(R)$ is a thin-plate spline function for R accounting for the display type
- $s_k^T(R)$ is a thin-plate spline function for R accounting for the comparison type

- D_j is the fixed effect of the display type (e.g. an intercept term)
- T_k is the fixed effect of the k^{th} comparison type
- P_l is the random effect of the l^{th} participant
- ϵ is the random error

Table 3.2: Parametric coefficients in gam model.

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.6862	0.1780	15.0931	0.0000
plot3dDigital	0.1357	0.1409	0.9626	0.3363
plot3dPrint	-0.0563	0.1415	-0.3980	0.6908
typeAdjacent	0.0968	0.1197	0.8088	0.4191

Table 3.3: Approximate significance of smooth terms in gam model.

Smooth	edf	Ref.df	F	p-value
s(Ratio):plot2dDigital	1.0005	1.0009	2.4500	0.1183
s(Ratio):plot3dDigital	1.7934	2.2056	0.7146	0.4521
s(Ratio):plot3dPrint	0.0003	0.0005	0.0004	0.9996
s(Ratio):typeSeparated	4.2371	4.7338	4.2693	0.0018
s(Ratio):typeAdjacent	1.0007	1.0014	0.0059	0.9429
s(participant)	27.0449	34.0000	4.7272	0.0000

A generalized additive model with random effects for subject was fit using the `mgcv` package (S. N. Wood 2004, 2003; S. N. Wood 2017). Spline terms by comparison type and display

type were included for the ratio variable, to account for indications of nonlinearity from Figure 3.11; cumulative predictions for fixed effects are shown in Figure 3.12. It is clear that there are differences between the errors for separated comparisons and adjacent comparisons; adjacent comparisons seem to have roughly similar errors for all ratio values, while separated comparisons seem to have lower errors for points near 0, .5, and 1, with higher errors as the distance increases from one of these anchor points; this effect has been explored in Hollands and Dyre (2000). Residuals from the model indicate that there is some left skew in the tails, but this is not extreme; code for this analysis can be found in the supplemental materials on GitHub.




Figure 3.12: Predictions with standard errors for fixed effects in a generalized additive model with splines for ratio by comparison and ratio by display method. It is clear that the separated comparisons are easier to estimate when small, right around 50%, or large, and harder to estimate between these points. What is interesting is that no such trend is present for adjacent comparisons.

While we were initially concerned that this effect was due to the slider initial setting of 0.5, we did examine this hypothesis; only 42 responses (8.81%) had slider values at exactly 0.5, and all but 6 were for ratios on either side of 0.5 (.464, .562). This suggests that while a few participants did take the strategy of using the default value for the slider, they did so strategically, not moving the slider only when the ratio was perceived to be sufficiently

close to 0.5. In addition, trials where participants utilized this strategy were evenly split between separated and adjacent comparisons, suggesting that this anchoring process is not the reason for the difference in spline fits for the different comparison types. It seems more likely that this effect is at least partially due instead to lower errors around the extremes of the ratio spectrum. A follow-up experiment with more ratios and in particular ratio values closer to suspected anchor points could explore this relationship in more depth - with only 7 different ratios, it is difficult to do more than note the appearance of an effect and speculate about the anchor points, particularly given that we have used spline models with 5 knots to fit the smooth.

Heer and Bostock (2010b) provided a zip file containing data and code for their paper, but this link has not been updated and is no longer active, so it is not possible to fit a similar model to their data for comparison purposes; if we could get access to their data, this would allow for investigating whether the differences identified in our model were an artifact of the measurement method (e.g. slider vs. numerical entry) or if this effect was present in previous studies and can now be identified because of more advanced statistical modeling techniques.

3.3.3 Interactivity with 3D Rendered Charts

When 3D rendered plots were shown, we recorded the number of user interactions (clicks, rotations) with the WebGL plot. **?@fig-clicks** shows the distribution of number of clicks over all 3D rendered trials in the experiment, as well as the number of clicks relative to the accuracy measure $\log_2(|\text{Judged Percent} - \text{True Percent}| + 1/8)$. It does not appear that participants who interacted with the WebGL rendering were more accurate than their peers

who did not.




Figure 3.13: Participant clicks on the WebGL interface (left) in order to interact with the plot. Most participants did not interact with the WebGL interface, perhaps because they did not realize that interacting was an option, but some participants utilized the interactivity heavily. (Right) Participant interactivity was not associated with increased accuracy.

3.4 Discussion and Future Work

Previous work in 3D graphics would suggest that the errors for the 3D graphs would be larger than the errors for the 2D graphs. While we did not find any significant results indicating that 3D graphs are read less accurately, there are two possibilities that might account for this discrepancy.

The first potential explanation is that this study is underpowered - the effect size is small, and our 38 participants were insufficient; the original study included 51 participants, which




Figure 3.14: Participant clicks on the WebGL interface (left) in order to interact with the plot. Most participants did not interact with the WebGL interface, perhaps because they did not realize that interacting was an option, but some participants utilized the interactivity heavily. (Right) Participant interactivity was not associated with increased accuracy.

is slightly larger. However, we should note that the analysis methods used in Heer and Bostock (2010b) and Cleveland and McGill (1984b) do not provide numerical tests of statistical significance, instead defaulting to graphical displays which include confidence intervals.

The second possibility is more interesting: we examined 3D charts using rendered 3D graphs and 3D printed charts; both of these options allow for participants to interact with the chart, rotating it, and generally perceiving it as one might perceive any other 3D, real, object. This is a far cry from the 3D perspective charts in the original study, which have a fixed angle and perspective and are thus not equivalent to our 3D charts. Future studies should include an additional fixed 3D perspective bar chart, which will at least enable us to examine whether modern 3D rendering environments allow for more accurate conclusions than fixed 3D perspectives. Future iterations of this study will include “traditional” 3D graphs created by Microsoft Excel (that is, graphs with a fixed 3D perspective rendered in 2D). This option will allow us to examine fixed perspective 3D plots compared to 3D renderings and 2D plots; it will also enable online data collection in addition to the in-person data collection used in this experiment.

Another interesting aspect of our study is that the method used to record participant estimates is different from the method used in the original study as well as Heer & Bostock’s replication study. A method similar to our slider input, marking position on a line, was used in Spence (1990), but Spence asked participants to estimate $A/(A + B)$, where we asked participants to estimate A/B ; thus, our results are still not directly comparable to previous studies. We expect that the specific ratio estimated would also have an effect on observed participant errors.

The slider method for input of ratio estimates should be easier for participants, as it does

not require explicit transformation to the numerical domain. What is clear is that it would be beneficial to assess the impact of the measurement method on participant errors directly, so that the results of these different studies might be explained and interpreted with regard both to the stimuli used and the measurement method employed in the experiment. Future iterations of this experiment will likely address this estimation difference; such modifications in experimental design are relatively straightforward in Shiny and will provide useful insight into the design of future experiments evaluating the perception of statistical graphics.

3.4.1 Supplemental Material

Stimuli, code, and data for this experiment are provided at <https://github.com/TWiedRW/2023-JDS-3dcharts>.

4 CHAPTER 2

4.1 Abstract

A key component of statistics courses is to teach students how to interpret data visualizations. Although many research-based recommendations exist for creating graphs, the technological advances for creating such graphs have outpaced studies that evaluate their effectiveness, especially when considering 3D graphs. Here, we describe a process for integrating an experiment on 3D graphs as part of a project on statistical investigations for students enrolled in an introductory statistics course and gathered responses as students reflected on their experience of being a participant in the experiment and then being a reviewer of empirical evidence about the experiment. A total of 82 students participated in our graphics project. As participants in an experiment, they displayed a pattern of not fully grasping research objectives as evidenced by widely varied responses. However, as reviewers of material from a pilot study of the same design, they tended to gain a clearer understanding of the purpose of the experiment and its role in the realm of data visualizations by correctly interpreting an extended abstract and video presentation of the pilot study. The project we presented to students shows promise as an educational tool for helping students gain a more holistic view of statistical research, which is important in both the contributions

to data visualizations and the education of statistics.

4.2 Introduction

The education of statistics emphasizes the use of real data and its application in answering research questions. Students in introductory statistics courses are mainly exposed to elementary methods and textbook examples that demonstrate the application of these methods, which is used, in part, because of the emphasis on teaching students to think statistically using real data (Carver, College, and Everson 2016). Many textbooks, such as Tintle et al. (2021), use scenarios and ask students to perform the corresponding inferential test. This process is then repeated over the course material without much deviation in the instructional method. In some cases, however, students can participate and benefit from well-designed classroom experiments (McGowan 2011). This process can expose students to concepts such as randomization and let students see the specifics of experimental design through their participation. Loy (2021) has demonstrated that student participants often recalled their experiment in later concepts, showing some evidence that students can benefit from the hands-on experience.

In addition to using experiential learning broadly, a key aspect in the statistics classroom is teaching students to interpret data through graphs. Nearly 40 years ago, Cleveland and McGill (1984a) began the process of establishing good practices for making graphs. While their findings have been replicated (Heer and Bostock 2010a), many areas in data visualization remain underdeveloped that can benefit from the framework used by Cleveland and McGill (1984a). One such area is the projections of 3D data visualizations. The current mantra is to avoid 3D graphs when possible and studies around the 1990s seem to

provide some valid skepticism of their use (Barfield and Robless 1989a; Fisher, Dempsey, and Marousky 1997a). Barfield and Robless (1989a) showed that participants were sometimes more accurate using 3D graphs than 2D graphs, depending on the participant’s experience level. However, participants were generally more confident with their answers for 2D graphs. Fisher, Dempsey, and Marousky (1997a) observed a similar preference for 2D graphs over 3D graphs when extracting information, but found no preference for visual appeal for the type of graph. While these results provided valid skepticism toward the use of 3D graphs, the results are generalized to the projections of the 3D graphs. The area of “true” 3D graphs is largely unexplored, but advances in technology allow for easier access to explore the 3D projections of these graphs. Kraus et al. (2020) explored the 3D projections with the use of virtual reality and found that participants were more accurate at extracting values from 3D heatmaps at the expense of needing more time than 2D alternatives.

Because the use of “true” 3D graphs remains largely unexplored, it provides a unique opportunity to be used as an experiential learning opportunity for statistics students. Not only can students benefit from the exposure to different graph types, including “true” 3D graphs, but they can also experience a more authentic method of teaching, which is more likely to be beneficial to reinforce statistical ways of thinking.

In this paper, we used an experiential learning project that employed different graph types, including “true” 3D graphs, in an introductory statistics classroom environment and described its potential application as an educational tool. We presented students with a series of modules where they first participated in our experiment on different graph projections, followed by acquiring knowledge on the purpose of the experiment by reading an extended abstract and watching a presentation of a pilot study of the same experiment.

4.3 Methods

4.3.1 Overview and Motivation

We introduced students enrolled in Introduction to Statistics (STAT 218) at the University of Nebraska-Lincoln to a graphics project that contained an experiment and progressively revealed components that illustrate the experiment’s research objectives. The project started by providing students with minimal information about the research objectives before revealing the scope of the experiment through an extended abstract and presentation. The goal of the graphics project was to observe how students think statistically about experiments as both participants and research consumers. Students were provided with mostly open-ended prompts throughout the graphics project, which allowed us to freely explore common themes in student responses without the limitation of preset choices.

4.3.2 Experiential Learning

The graphics project was split into two main components regarding the interaction of students with the material. In this section, we will discuss student interactions with the experiment. Students were first presented with a series of four modules presented from the role of research participation. These modules contained the informed consent form, pre-experiment reflection, experiment participation, and post-experiment reflection. Within the informed consent module, students were informed that their data would be anonymized and that the experiment was exempt from the institutional review board (Project ID: 22579). While all students were given the option to participate in the graphics project, we were only able to collect responses when informed consent was obtained and if the student was

19 years of age or older. In the pre-experiment reflection, students were asked to write a paragraph about how they think the process of scientific investigation looks from the perspective of researchers and the general public. The experiment participation module asked students to paste the code generated upon completion of the graphics experiment, which is detailed in the next section. The generated code serves as a basic check to indicate whether or not students fully participated in the experiment. For the post-experiment reflection, students were asked five questions about the purpose of the experiment. These include questions on the hypotheses being tested, sources of error, variables of interest, and elements of experimental design.

After completing the experiment reflections, students moved to the reflection of the overall research objectives. Students were first directed to read an unpublished two-page extended abstract that we submitted as a contributed paper for the Symposium on Data Science & Statistics (SDSS). The extended abstract outlined the experiment’s purpose and procedures, but not the results from our initial pilot study. After reading the extended abstract, students were asked to write a paragraph about what they found clearer about the experiment’s purpose than when they were a participant. Finally, students were directed to watch a 12-minute pre-recorded presentation based on an abbreviated version given at SDSS (Wiederich 2023). The video contained the same material as the extended abstract and included the results from our pilot study. The presentation reflection asked students four questions about the experiment and how information was presented differently than in the extended abstract.

Except for the informed consent and experiment participation modules, all student responses were open-ended. Each question and its corresponding module can be found in Table 1.

Instructors for STAT 218 were recruited for Summer 2023 and Fall 2023 to administer the graphics project into their classroom. The instructors were given the option of administering the project as coursework material or extra credit, along with the liberty of grading at their discretion. Instructors for the June-July Summer 2023 courses did not have the abstract or presentation modules.

4.3.3 Data Analysis

Since nearly all of the student responses to the project modules are open-ended, the analysis of the project modules is qualitative. We will selectively extract student responses that demonstrate variability and repetitive themes in their understanding of the graphics experiment. For all modules, except the Post-Experiment reflection, students are asked to comment on the experiment without an objectively correct response.

For paragraph responses in the Pre-Experiment and Abstract reflections, bigram plots are used as a graphical analysis to display word pairs after removing stop words (e.g., “the” and “and”). These word pairs help illustrate common themes that students wrote about in their longer prompts. In the Post-Experiment reflection, the prompts have objectively correct answers but would require a subjective aggregation to indicate the level of correctness. Instead, we opt to select at least one student response that we consider to be “most correct” and a couple of other responses that are either partially correct or entirely incorrect to illustrate the variability of the responses. The results of the graphics experiment are presented by Wiederich and VanderPlas (n.d.) as an extension of a larger study comparing the dimensionality and projections of 2D and 3D bar charts.

4.3.4 Graphics Experiment

4.3.4.1 Constructing Stimuli

Based on Cleveland and McGill’s seminal work (1984a) on graphical perception, participants were presented with a series of bar graphs where two bars are marked with either a circle or a triangle. The heights of the bars were chosen from the following equation:

$$s_i = 10 \cdot 10^{(i-1)/12}, \quad i = 1, 2, 3, \dots, 10 \quad (4.1)$$

where s_i represents a value given an integer i as defined above. The values s_i from Equation 4.1 were then paired such that the ratio of the smaller value to the larger value yield the ratios of 0.178, 0.261, 0.383, 0.464, 0.562, 0.681, and 0.825. Each bar graph has two groupings of five bars. Following the Type 1 and Type 3 graphs from the position-length experiment by Cleveland and McGill (1984a), the value pairs for each ratio were either placed in the first grouping on the second and third bars (adjacent) or placed on the second bars in each grouping (separated).

To explore the effect of dimensionality and projection of the bar graphs, we introduced the following plot types: 2D digital, 3D digital (static), 3D digital (interactive), and 3D printed. There was no single software package that could create all four plot types, so we carefully constructed graphs from different software packages to be as similar as possible. The 2D digital plots were rendered with the `ggplot2` package (Wickham 2016). Microsoft Excel® was used to render the 3D digital (static) plots. The 3D digital (interactive) and 3D printed plots were created with OpenSCAD® (Kintel 2023), where the generated stereolithography (STL) files for the 3D digital (interactive) plots were rendered with the `rgl`

package (Murdoch and Adler 2023).

4.3.4.2 Experimental Design

With 56 treatment combinations, we opted to use an incomplete block design to provide participants with 15-20 graphs. Kits of graphs were constructed so that five of the seven ratios are equally represented, resulting in 21 different kits. Within each kit, all graph types appeared for each ratio and the comparison type was randomly assigned. A visual layout of the experiment is shown in Figure 1. All kits received a unique identifier and a set of instructions for accessing the experiment.

A Shiny application (Chang et al. 2024) was designed to administer the experiment. Students were directed to randomly select a kit of graphs and visit the Shiny application's website linked on the instructions. For students enrolled in the online sections of STAT 218, the website was provided by the instructor and they were prompted in the application to select that they were an online participant; selecting online participation resulted in the 3D-printed plots being removed from the set of graphs presented to the participant. After students provided a kit identifier (if applicable), students were presented with graphs in a randomized order. If the student marked that they were an online participant, the 3D-printed graphs were removed from their experiment lineup. Each graph asked the students to first identify the larger marked bar and then to guess the height of the smaller marked bar if the larger marked bar was 100 units tall using a slider widget. After completing the experiment, a completion code was generated for students to paste into the experiment participation module.

4.4 Results

4.4.1 Recruitment of Students and Instructors

We recruited 3 instructors for the summer and fall semesters in 2023. Each instructor offered the project as extra credit in their course and student participation was entirely voluntary. A total of 82 students participated in the project, and 9 students did not complete the entire project (Table 2).

4.4.2 Selected Responses from Experiment Participation

4.4.2.1 Pre-Experiment Reflection

In the first stage of the project, before participating in the experiment, students had limited information about the research objectives and we recorded how students thought about scientific research. Students were initially prompted with one question, “What do you think [the] process [of scientific investigation] looks like, from the perspective of a researcher, compared to what it looks like from the perspective of someone in the general public who is a consumer of scientific results?” The responses to this question allowed us to gather a broad understanding of what students initially thought about the research process. Students generally understood the purpose of scientific research by connecting the ideas of hypothesis testing and publishing results. Students wrote about scientific research starting from the place of a question, followed by conducting an experiment and relaying the results to the public. For example, the most common phrase groupings include variations of “scientific research”, “data collection”, and “peer review”. A bigram plot of student responses to the Pre-Experiment reflection prompt is shown in Figure 2.

4.4.2.2 Post-Experiment Reflection

After participating in the experiment, students were provided prompts asking about the goals of the research objectives, which were to evaluate the errors of ratio judgments between 2D and 3D graph types under different comparison conditions in a randomized block design. Some students correctly identified parts of the questions asked in the post-experiment reflection, but most students typically missed some or most of what would be considered a correct answer.

We first asked students “What do you think the purpose of the experiment was?” For this question, a complete response would have included the comparisons of ratio judgments of the projections for 2D and 3D graphs. A student example of the correct response was “I think this experiment aimed to test if it was easier to compare two graphs in 2D or 3D.” Another student missed the comparison of 2D to 3D graphs, but was mostly correct with their response “I think the purpose of this experiment was for the researcher to gather data on how people perceive, interpret, and understand 3D graphs.” One student incorrectly thought that we were testing differences in demographics by responding “They could be trying to determine how different genders, ages, etc. perceive the sizes of the bars in the graph. Demographics could make a pretty significant difference.”

We then asked students “What hypotheses might the experimenters have been testing?” A correct response would include testing the differences in errors of ratio judgments between graph types. One example of a student providing a correct response was given by stating “They might have been testing if a 2D model is easier to estimate its relative size to another when compared to a 3D model of it.” Other students replied with statements that would not be able to be measured from the experiment, with one student responding “How taking

Statistics 218 effects how you can compare two groups” and another student saying “That 2d is preferred over 3d. It cleans up the data presentation.”

An important topic covered in STAT 218 is randomization, which we asked students with the prompt “What elements of experimental design, such as randomization or the use of a control group, do you think were present in the experiment? Why?” In the experiment, randomization was performed with the ordering and assignment of graphs into kits. One student perfectly described randomization by responding with “Randomization: The survey used an experimental design where in the survey there were different sets of 3D charts and maybe by a randomization process each participant was shown a different set of charts to see the differences in interpretations of the charts based on which set was assigned. Control Group: Since this survey aimed to only understand how participants interpret 3D charts without comparison to other chart types, then no control group was needed.” Another student was partially correct with a response of “Randomization was used because the ever person got a different graph.” Other students missed the use of randomization, with one student responding “random students in the stats class” and another saying “Randomization was not used because it was offered as an extra credit assignment in class.”

4.4.3 Selected Responses from Research Reflections

In this set of reflections, students first read the extended abstract, followed by watching the 12-minute presentation video. The abstract unveiled the scope of the study to students, many of whom did not realize the underlying complexities. Nearly all students responded with statements about gaining clarity about the purpose of the experiment and its role in testing the differences between 2D and 3D graphs. Students commented on how they now

understood the purpose of comparing 2D and 3D graphs, and also the potential benefits that may stem from research, such as graphical accessibility to the visually impaired. A bigram plot of the student responses to the abstract reflection prompt is shown in Figure 3.

Lastly, more than half of the students (78.5%) responded that they preferred the presentation over the extended abstract when asked “If you had to hear about this study using only the extended abstract or only the presentation, which one would you prefer? Which one would be better for determining whether the experiment was well designed?” One student responded “I am a visual learner so I would have rather heard about in through the presentation. It also broke down the steps which is easier for me to understand. I think the presentation as a whole would be better for determining how the experiment is designed.” Another student said “Personally I like the abstract better. If I get confused on something it is so much easier to go back and reread to understand what is going on. If I ask myself questions about it, it is much easier to go back and find answers to the questions as well.”

4.5 Discussion

Taken together, our results support the idea that we achieved our goal of providing students of an introductory statistics course with the opportunity to reflect on active research. Students generally appreciated the progressively revealing nature of the graphics project, which is evident from the abstract and presentation reflections. When provided with the post-experiment reflections, students often either missed the research objectives of the experiment or had partially correct responses. However, this was expected since the design of the experiment exceeds the scope of experimental design taught in typical introductory

statistics courses.

The abstract reflection received many responses indicating that students had moments of realization about the true nature of our research goals, which was further expanded with the presentation reflection prompts. Across all reflections, students were thoughtful, and sometimes amusing, with their responses and they were on the path of statistical thinking.

Having students participate in experiments is not uncommon, possibly attributed to the readily available convenient samples within academia and potential applications to introduce new course material (Margaret 1994). However, these studies typically end for students after completing the experiment (McGowan 2011; Zacks et al. 1998; Fisher, Dempsey, and Marousky 1997a). In our study, we integrated the graphics experiment as a course project with the addition of reviewing research material.

What we found was that the graphics project found success within the recommendations of the GAISE College Report by using real data to illustrate the contextual purpose of scientific research (Carver, College, and Everson 2016). Students were able to demonstrate their ability to think statistically through the series of reflections. The graphics project made use of real data, along with data collection, within the scope of an approachable and field-related topic, which allowed students to see how scientific research is conducted in the field of statistics.

A limitation of our study is the use of open-ended responses that do not objectively assess student learning. While the student responses were useful in gathering insight, the responses are widely varied and do not have direct comparisons of statistical thinking throughout the modules. Another limiting factor is the tiered layering of convenience sampling, with instructors being recruited before recruiting students, which impacts the generalization of

our findings. Nonetheless, the 82 students who participated provided meaningful answers that displayed some level of statistical thinking throughout the graphics project.

Future studies could use a similar framework to conduct experiments on more typical 3-dimensional data, such as heatmaps. The use of graphical experiments in the classroom not only provides a readily available convenience sample but also adheres to the recommendations of the GAISE College Report (Carver, College, and Everson 2016). With the framework we provided in this paper, we aim to make adjustments to further improve the graphics experiment and corresponding project as an experiential learning opportunity.

4.6 Figure legends

Figure 1: Visual display of the experimental design for students who participated in the 3D bar charts experiment. Kits of graphs were created by first choosing five ratios from nine available options (1). Each ratio then uses all graph types, with the exception of the 3D-printed graphs for online students (2). Finally, all graphs were randomly assigned to have the marked bars as adjacent or separated (3).

Figure 2: Bigram of student responses (n=82) to the pre-experiment prompt. Each line represents pairs of words that appeared together where each pair occurred at least twice. Students generally understood that science is about investigating research questions and collecting data.

Figure 3: Bigram of student responses (n=63) to the abstract reflection prompt. Each line represents pairs of words that appeared together where each pair occurred at least twice. Students generally focused on the research objective of comparing 2D and 3D graphs.

4.7 Tables

Table 4.1: Questions provided to students in each project module.

Reflection	Question	Prompt
Pre-Experiment	Q3	In this class, you'll be learning about the process of scientific investigation. What do you think that process looks like, from the perspective of a researcher, compared to what it looks like from the perspective of someone in the general public who is a consumer of scientific results? Write a paragraph (at least 3-5 sentences) about how you think science happens.
Post-Experiment	Q5	What do you think the purpose of the experiment was?
	Q6	What hypotheses might the experimenter have been testing?
	Q7	What sources of error are involved in this experiment?
	Q8	What variables were examined? For each variable, identify whether it was quantitative or categorical.
	Q9	What elements of experimental design, such as randomization or the use of a control group, do you think were present in the experiment? Why?
Abstract	Q10	What components of the experiment are clearer now than they were as a participant? What questions do you still have for the experimenter? Write 3-5 sentences reflecting on the abstract.

Presentation	Q11	How did the information you gained from the components of this project (participation, post-study reflection, extended abstract, presentation) differ?
	Q12	What components were emphasized in the presentation that weren't emphasized in the abstract? Why do you think that is?
	Q13	What critiques do you have of this study and its design? What would have made the study better?
	Q14	If you had to hear about this study using only the extended abstract or only the presentation, which one would you prefer? Which one would be better for determining whether the experiment was well designed?

Table 4.2: Number of valid student participants by semester.

Semester	Number of Sections	Number of Students
Summer 2023 (May-June)	1	17
Summer 2023 (July-Aug)	1	23
Fall 2023 (May-June)	1	42

Students under 19 years of age or did not consent were excluded from data collection. To comply with IRB, no demographic information was collected to keep students anonymous with their reflections.

5 CHAPTER 3

5.1 Introduction

In the 20th century, advancements in technology made data visualizations increasingly more affordable and accessible to a broader population. The primary change in formal data visualizations was from hand-drawn charts to computer-rendered charts (Tukey 1965), yet other technological advances have allowed for data visualizations to enter other mediums. These include the ability to effectively use the 3-dimensional (3D) world around us with the novel use of 3D-printed charts. This newer type of visualization has gained a small amount of traction in recent years as a method of producing tangible charts.

At this time, there are few, if any, studies that evaluate 3D-printed charts for the purpose of displaying statistical information. This may in part be due to the cost of materials and rendering times as compared to charts produced on a computer. A single chart can take up to a day to print, limiting the ability to quickly produce visualizations. Given the nature of 3-dimensional data, viewing a 3D realization of statistical information in a 3D environment is a realistic scenario that could increase the ability for information extraction.

3D-printed visualizations have shown mixed or promising results across many disciplines.

Katsioloudis and Jones (2018) showed no evidence of a statistical difference in the method of 3D renderings when tasked with drawing a cross-sectional of a dodecahedron. This demonstrated that spatial awareness between computer-rendered and 3D-printed shapes is not largely different among engineering students. The use of 3D-printed maps for navigation for visually impaired persons showed positive feedback by Holloway et al. (2019), increasing the accessibility of navigation. In the clinical setting, 3D-printed anatomy structures were well accepted along with VR-glasses and 3D computer renderings (Muff et al. 2022). With the rise of 3D-printed visualizations in scientific fields, we turn our attention to their use for statistical graphics.

5.1.1 Literature Review

An identical dataset across multiple chart types does not ensure that the data is perceived in the same way (Cleveland and McGill 1984a; Hofmann et al. 2012; Vanderplas, Cook, and Hofmann 2020b). One of the main factors to this phenomena is with how data is encoded into the chart. These encodings include, but are not limited to, placement along axes, lengths, areas, volumes, and color scales. For example, bar charts and pie charts are two common visualizations that have long since been the topic of debate (Eells 1926b; Croxton and Stryker 1927a; Cleveland and McGill 1984a). In the case of bar charts and pie charts, encodings are represented by lengths and angles, respectively. Inherently, the comparison of different chart types is a comparison of encodings due to the changes in how data is being displayed.

Cleveland and McGill (1984a) noted that estimates involving numerical accuracy may decrease when increasing dimensionality of the encoding, although this was not formally tested

in their experiments. The reasoning is possibly due to Stevens' power law, a mathematical formulation of how magnitudes are perceived given different stimuli sources (Stevens and Stevens 1986). The general form of the law is $\psi(I) = kI^\alpha$, where I is the magnitude of a stimulus, $\psi(I)$ is the perceived magnitude, k is a proportionality constant from the unit of the stimulus, and α is the exponent from the type of stimuli used. Studies have estimated that lengths are perceived without bias (i.e., $\alpha = 1$), but areas and volumes tend to have skewed perceptions (Cleveland and McGill 1984a). This indicates that lower-dimensional charts might perform better when readers are tasked with extracting numerical estimates from the chart.

There are mixed results in regard to the use of 3D charts, mostly attributing to the purpose of third dimension. When the extra dimension does not convey meaningful information, estimates of accuracy decrease and solution times increase as compared to equivalent 2D charts (Fisher, Dempsey, and Marousky 1997a; Zacks et al. 1998; Fischer 2000). The same increase in solution time is seen when the third dimension is utilized for displaying data, but can sometimes produce better error rates than 2D charts (Barfield and Robless 1989a; Kraus et al. 2020). Additionally, when given the option of 2D or 3D charts for extracting numerical information, the 2D charts showed increased preference and confidence than their 3D counterparts (Barfield and Robless 1989a; Fisher, Dempsey, and Marousky 1997a). It is worth noting that all of the studies listed use renderings of 3D charts and not physical 3D charts.

Formal studies involving true 3D charts are limited, and it is unclear if they follow the framework of existing theories on data visualizations. Unlike paper and computer rendered charts, constructing true 3D charts inherently contains many additional factors that could affect perception, such as chart materials, natural lighting, interactivity, and viewing distance.

Some of these factors have already been shown to have an effect on computer renderings (Tarr and Kriegman 2001; Wang et al. 2022).

We hypothesize that 3D charts in 3D environments will produce better information extraction than their computer rendered counterparts. Specifically, we will compare 3D-printed charts to digital 2D and 3D renderings. We suspect that this difference will hold across multiple data sets and different magnitudes of stimuli. In this paper, we evaluate the accuracy of numerical estimations on true 3D charts by conducting a factorial experiment that assessed chart types and ratios of pairs of stimuli. We discuss the construction of the stimuli and how we closely matched the charts to compare 2D, 3D-digital, and 3D-printed renderings of heatmap data.

5.2 Methods

Our study is designed to evaluate and expand the literature on numerical estimation of 3D charts. In our study, our focus is on 2D and 3D heatmaps, which we carefully construct to ensure that differences are contributed to the dimensionality of the chart. All of our data and methods are publically available for open science and reproducibility at <https://github.com/TWiedRW/ch3-heat3d>. In this section, we discuss process of designing our experiment and participant recruitment.

5.2.1 Stimuli

We denote “stimuli” to represent the magnitude of our chosen values. In a 3D Cartesian space, the X- and Y- axes represent the coordinates of the stimuli, and the Z-axis represents the value for the stimuli. Each X and Y coordinate is represented by a square tile with

a 1:1 aspect ratio. For a 2D space, the Z-axis is replaced by a color gradient scale. All stimuli and remaining randomly generated values range between 0 and 100 units. In this experiment, $X = 1, 2, \dots, 10$ and $Y = 1, 2, \dots, 10$.

The design of our experiment made use of the method of constant stimuli, where comparisons between stimuli are with respect to a stimuli that remains the same magnitude (Kingdom and Prins 2016, chap. 3). We set the constant stimuli at 50 units. For stimuli between 50 and 100, we set the maximum stimuli value at 90 and equally partitioned the ratios of stimuli with the constant stimuli, $\frac{50}{90} = 0.556$ to $\frac{50}{50} = 1.0$, resulting in 4 varying stimuli values where 50 is the smallest value. The same ratios obtained with stimuli between 50 and 90 were used to create 4 stimuli varying between 0 and 50, where 50 is the largest value. Additionally, we also included a stimuli pair where both values are 50, resulting in nine total pairs of stimuli. All stimuli values can be found in Figure 5.1.




Figure 5.1: Values for stimuli in the heatmap experiment. All values are paired with the constant stimuli of 50 units, creating nine pairs of stimuli.

To generate non-stimuli random values, we used a mixture distribution of random uniform noise and a mathematical function to populate our coordinate grid. The mathemat-

ical functions are scaled between 0 and 100, $g(Z) = 100 \cdot \frac{Z - \min(Z)}{\max(Z) - \min(Z)}$. Two datasets were created for the experiment. The first dataset, called set 1, used the formula for the top half of sphere that is centered within our X and Y coordinate grid, $f_1(X, Y) = \sqrt{7^2 - (X - \bar{X})^2 - (Y - \bar{Y})^2}$, where $\bar{X}$ and $\bar{Y}$ are the averages of the X and Y coordinate ranges. The second dataset is calculated similarly using the formula for the bottom half of sphere, $f_2(X, Y) = \sqrt{7^2 - (X - \bar{X})^2 + (Y - \bar{Y})^2}$. Denoting Z as the random values, $U(0, 100)$ as a random variable drawn from a continuous uniform distribution, and X, Y as heatmap coordinates, our random heatmap data is calculated in Equation 5.1 with $c = 0.3$. An example of varying c values is presented in Figure 5.2.

$$Z = c \cdot U(0, 100) + (1 - c) \cdot g(f_i(X, Y)) \quad (5.1)$$




Figure 5.2: Mixture distribution of Equation 5.1 using the formula for the top half of a sphere. As c approaches 1, the distribution resembles uniform random noise.

The placement of stimuli values onto the randomly generated heatmap data was done via simulation to try to make their placement look as natural as possible. Twenty random heatmaps were generated. For each heatmap, the non-constant stimuli was placed onto the coordinate such that the difference between the stimuli and randomly generated value is minimized. The constant stimuli was then placed onto a coordinate with a Manhattan distance of three or four that minimizes the difference between the constant stimuli and

randomly generated values, where the Manhattan distance is given by $|X_i - X_j| + |Y_i - Y_j|$ for stimuli i and j . To ensure that stimuli placement is evenly position across the heatmap, the count of stimuli was computed separately across the X and Y axes. For example, in Data Set 1, the X-axis has four stimuli in $X = 1$, three stimuli in $X = 2$, and so forth. Chi-squared statistics were calculated for each axis and the heatmap with the smallest average Chi-squared statistic was selected as the final dataset. A visual inspection of this process showed that the stimuli were not clustered in any one area of the chart and that the stimuli look natural with respect to the random mixture distribution.

5.2.2 Charts

Three types of charts were considered for this study: 2D-digital (2dd), 3D-digital (3dd), and 3D-printed (3dp). We constructed these charts so that they are as similar as possible, but inherent difference between dimensionality led to artistic decisions that attempt to focus solely on the dimensionality of the charts. The process of creating the charts is discussed in this section.

The 3D-printed charts were rendered with OpenSCAD (Kintel 2023). To include plot text, a 120mm by 120mm by 10mm base was created with a solid color that was either white or black. Cells of the heatmap measured 10mm by 10mm, resulting in a heatmap that is 100mm by 100mm and is centered on the base. The upper bound of the height of heatmap values is 100mm, where 1-unit in the heatmap data is represented by 1mm of height on the heatmap. Once rendered, the heatmap was saved to 3D Manufacturing Format (3mf) and Standard Triangle Language (stl) files. A variety of solid and gradient filaments were used to print the output files from OpenSCAD. An example of the 3D-printed chart is shown in

Figure 5.3.



(a) Solid filament of Data Set 1

(b) Solid filament of Data Set 2

Figure 5.3: 3D printed heatmaps.

To closely match the 3D-digital chart to the 3D-printed chart, multiple stl files were created for each colored component and combined with the RGL package (Murdoch and Adler 2023). The base was rendered with white smoke (`#F5F5F5`) to slightly contrast with the default white background color (`#FFFFFF`). Heatmap tiles were rendered with cyan (`#74CCFF`) and text labels were rendered with black (`#000000`). Lighting was fixed at 45 degree angles at two opposite corners of the chart. The end result was a near perfect replica of the 3D-printed charts, with the exception of different heatmap tile colors, where an example is given in Figure 5.4.

Unlike the 3D charts, the 2D charts needed a different encoding to convey heatmap values. The 2D heatmaps were created with `ggplot2` (Wickham 2016) using `geom_tile()`. Fill




Figure 5.4: RGL rendering of Data Set 2 for 3dd charts.

colors for the cells use a color gradient from Blue Zodiac (#0C2841) to Malibu (#66D9FF), which were selected from a color picker using shadows on our initial 3D-printed charts. The color interpolation was performed with the `scale_fill_gradient()` function from the `ggplot2` package.



Figure 5.5: Color palette for 2D-digital charts. The colors are interpolated from #0C2841 to #66D9FF, which are the the colors of the lighting conditions for a 3D-printed chart created with cyan filament.

5.2.3 Subject Recruitment

Participants were recruited from all STAT 218: *Introduction to Statistics* sections at University of Nebraska-Lincoln from June 2025 to Spring 2026. The experiment was incorporated into the curriculum as a project that allowed students to get hands-on experience with statistics. Throughout the project, students were also given a series of reflections to evaluate their perspectives on statistics. The first reflection had gathered students' thoughts on the scientific process before participating in the experiment. After the experiment participation, students reflected on what they thought the experiment was measuring. Lastly, students were given a written article and video presentation of the experiment and asked how their view of the experiment changed from their initial perspectives of the experiment's purpose.

To assess whether students completed the experiment or not, the experiment produced a completion code for students to submit to Canvas. The completion code was created by concatenating three words with dashes (e.g., palm-raising-creatively). Words for the completion code came from the movie transcripts of the Star Wars prequel movies (Episodes

1, 2, and 3) as presented by movies.fandom.com. The movie transcripts were tokenism and filtered to remove colorful language¹, stop words², and fictional words³. Instructors were provided an R script to check if their students completed the experiment, but the completion codes were not saved with any identifying information that could link experiment responses to student reflection responses.

5.2.4 Experimental Design

Our experiment was designed with a $3 \times 2 \times 9$ treatment structure. Media type is our main interest, with 2D-digital, 3D-digital, and 3D-printed charts. To ensure that results are not confounded with datasets, we used two datasets to create the heatmaps. A total of 9 stimuli pairs were placed into each dataset. The order of treatments was given so that media and dataset combinations were grouped together randomly in the sequence and stimuli pairs were randomized within the groupings.

Due to practical constraints, stimuli pairs were incompletely blocked. A full factorial design would result in 54 trials per participant, which could lead to a decrease in quality responses (Herzog and Bachman 1981). Therefore, we selected four out of the nine possible stimuli pairs to create incomplete blocks. This resulted in 18 blocks for a balanced incomplete block design. Within each block, media type is fully crossed with dataset. Using the incomplete block structure, participants completed a total of 24 trials, which is more practical than the full factorial design. Blocks were chosen for each participant by using probabilities that are proportional to the inverse counts of used blocks, meaning

We measured three responses – two questions for each trial and one question for each media

¹<https://www.cs.cmu.edu/~biglou/resources/bad-words.txt>

²tm package (Feinerer and Hornik 2024)

³Apple Inc. (n.d.)

by dataset combination. For each trial, participants are initially asked which stimuli in a pair is larger or if the stimuli are the same value. Next, they were asked to estimate the magnitude of the smaller stimuli if the larger stimuli represents 100 units, which is a subtractive process (Veit, n.d.). After each grouping of media and dataset, participants were presented with a modal dialog window asking them to rate their confidence on a 5-point Likert scale for their answers of the previous group of charts.

5.2.4.1 Shiny Application

A Shiny application (Chang et al. 2024) was developed to administer the experiment. The application consisted of five sections: informed consent, demographics, practice, experiment, and wrap-up. The entire application is designed to be completed in approximately 30 minutes.

The Shiny application started with the informed consent screen, allowing participants to select if they are a STAT 218 Student and if they agree to the data collection. Participants had to select a data collection option to continue. After submitting their data collection response, a completion code was generated and saved on the last page of the application. A copy of the informed consent is available in our GitHub repository.

If participants agreed to have their data collected, they were presented with the demographics section. This section asked participants to use drop down menus to specify their age, gender identity, highest education level, and a question about how their participation in the experiment is graded. The last question was a text box and asked participants to specify their favorite movie and/or actor. Demographic information was combined with the application start time and completion code to create a hash for an anonymized participant

identifier using the `rlang` package (Henry and Wickham 2025).

Once a participant completed the demographics page or selected “No” to the data collection question, participants were given a practice page. A modal dialog window was initially shown with instructions, and users could display this window again at any point during the practice. The practice consisted of four questions: two trials from 2dd charts and two trials from 3dd charts. Each practice trial showed the correct solution after the participant submits their trial. After all practice trials were completed, another modal dialog window was displayed to ask participants if they have access to the 3D-printed charts. This was necessary since the experiment was given to potentially online sections of STAT 218 who did not have access to the physical charts.

The experiment page was presented to participants after completion of the practice trials Figure 5.6. Each page contained two questions – one question for identifying which stimuli is larger and another question for estimating the value of the smaller stimuli if the larger stimuli is 100 units. If participants selected that the stimuli are the same value, then the slider was automatically placed at 100 units. Each trial has the slider initially placed at 50 units. An option was provided for displaying the coordinates of the stimuli pair on a 2D heatmap, but was left unchecked by default. For 3dd charts, the number of interactive clicks was also recorded. A trial could only be submitted if the first question is answered and if the slider was moved at least once.

After completing all trials, participants were given the completion code and informed to copy the code to Canvas since they will not have access to the code after closing out of the experiment.




Figure 5.6: Graphics experiment page. Each trial contained two questions about a pair of stimuli, and a check box that displays the location of the stimuli if needed.

5.2.4.2 Method of Analysis

The design of our experiment is a balanced incomplete split-plot with replication, where media type and dataset is the whole plot and stimuli pair is the split plot. A similar design was presented by Mandal, Parsad, and Dash (2020) for a single replicate. For our experiment, the ANOVA table is in Table 5.1.

Table 5.1: ANOVA table for analysis of a single replicate of the balanced incomplete split-plot.

Source	Degrees of Freedom
Block (B)	18-1
Media type (M)	3-1
Data Set (D)	2-1
Whole-plot error = $B \times M \times D$	$(18-1)(3-1)(2-1)$
Stimuli Pair (P)	9-1

Source	Degrees of Freedom
$P \times M$	$(9-1)(3-1)$
$P \times D$	$(9-1)(2-1)$
$P \times M \times D$	$(9-1)(3-1)(2-1)$
Split-plot error = $B \times M \times D \times P$	$(18-1)(3-1)(2-1)(9-1)$

The first question for each trial was to indicate which value within a stimuli pair was larger, or if they were the same value. Since a participant could either answer this question correctly or incorrectly, a Binomial generalized linear mixed model (Stroup et al., n.d.) was fit to the responses.

To measure numerical accuracy of participant responses for the second question in each trial, we used the absolute value of the difference between the participant estimate of the smaller stimuli value and its true value. By using the absolute value of the error, we measure how far off a participant's estimate was rather than the direction of error. For numerical accuracy, we fit a linear mixed model. If the responses do not fit well to a Gaussian distribution, other distributions, such as the log-normal or gamma distributions, will be used for a generalized linear mixed model.

$$Y = |\text{Guess} - \text{Actual}|$$

In addition to linear mixed models, we used the bootstrap distribution of mean errors as described by Cleveland and McGill (1984a) to provide empirical claims for the effect of dimensionality. A common issue in experiments involving perception is that random noise

often covers the effects of interest, so the empirical claims may provide additional insight into the effect of chart dimensionality (XXX).

After each media type and data set grouping, participants were given a question to rate their confidence on their answers. Here, we use descriptive statistics and visualizations to assess confidence across our treatment factors.

5.3 Limitations

Our study was designed to evaluate the accuracy of numerical estimation of stimuli pairs across three types of charts. However, numerical accuracy is not the only deciding factor in what makes a good chart design. Other factors include comprehension, recall, and overall preference (Fisher, Dempsey, and Marousky 1997a; Vanderplas, Cook, and Hofmann 2020b). Given the lack of current research into 3D-printed statistical graphics, we hope to provide some baseline insights into their use as a statistical tool.

We did not account for potential visual illusions in the construction of the heatmap data. A phenomenon seen in statistical graphics studies is that surrounding values may attribute to participant estimates (Zacks et al. 1998; VanderPlas and Hofmann 2015b). In our study, the placement of stimuli was so that the stimuli looked naturally positioned within the structure of the data, resembling patterns typically seen in spatial data. However, consistency of neighboring values was not accounted for in this experiment and would be worth having a follow up study to measure the effect of surrounding values on 3D-printed charts.

Lastly, we wanted to construct heatmaps so that they closely resembled each other across

dimensionality. The 3dd charts are identical to the 3dp charts, with the exception of heatmap colors. Here, the issue is that 3D heatmaps are not typically constructed by compiling stl files, but rather using packages such like `lattice` package (Sarkar 2008) to display the data. Due to parameter settings and code limitations, it may not be possible to get typical 3D renderings of charts to appear visually identical to 3D-printed counterparts.

5.4 Conclusion

Our study provides insight into the novel use of 3D-printed statistical graphics. We hope that these charts may be useful for educational purposes, accessibility of visually impaired persons, or recreational use of statistical communication.

6 REFERENCES

- 1978 *Handbook of Agricultural Charts. Enlargements* / [u.s. Department of Agriculture].
1978. District of Columbia: United States. Department of Agriculture; U.S. Dept. of Agriculture ; Distributed by Science; Education Administration.
- Apple Inc. n.d. “Web2 Dictionary File.” [/usr/share/dict/web2](#).
- Barfield, Woodrow, and Robert Robless. 1989b. “The Effects of Two- or Three-Dimensional Graphics on the Problem-Solving Performance of Experienced and Novice Decision Makers.” *Behaviour & Information Technology* 8 (5): 369–85. <https://doi.org/10.1080/01449298908914567>.
- . 1989a. “The Effects of Two- or Three-Dimensional Graphics on the Problem-Solving Performance of Experienced and Novice Decision Makers.” *Behaviour & Information Technology* 8 (5): 369–85. <https://doi.org/10.1080/01449298908914567>.
- Bateman, Scott, Regan L. Mandryk, Carl Gutwin, Aaron Genest, David McDine, and Christopher Brooks. 2010. “CHI ’10: CHI Conference on Human Factors in Computing Systems.” In, 2573–82. Atlanta Georgia USA: ACM. <https://doi.org/10.1145/1753326.1753716>.
- Borgo, R., A. Abdul-Rahman, F. Mohamed, P. W. Grant, I. Reppa, L. Floridi, and Min Chen. 2012. “An Empirical Study on Using Visual Embellishments in Visualization.”

- IEEE Transactions on Visualization and Computer Graphics* 18 (12): 2759–68. <https://doi.org/10.1109/TVCG.2012.197>.
- Brath, Richard. 2014. “3D InfoVis Is Here to Stay: Deal with It.” In *2014 IEEE VIS International Workshop on 3DVis (3DVis)*, 25–31. Paris, France: IEEE. <https://doi.org/10.1109/3DVis.2014.7160096>.
- Breslow, Leonard A., Raj M. Ratwani, and J. Gregory Trafton. 2009. “Cognitive Models of the Influence of Color Scale on Data Visualization Tasks.” *Human Factors: The Journal of the Human Factors and Ergonomics Society* 51 (3): 321–38. <https://doi.org/10.1177/0018720809338286>.
- Carver, Robert, Stonehill College, and Michelle Everson. 2016. “Guidelines for Assessment and Instruction in Statistics Education (GAISE) College Report.”
- Chang, Winston, Joe Cheng, JJ Allaire, Carson Sievert, Barret Schloerke, Yihui Xie, Jeff Allen, Jonathan McPherson, Alan Dipert, and Barbara Borges. 2024. “Shiny: Web Application Framework for r.” <https://shiny.posit.co/>.
- Cleveland, William S., and Robert McGill. 1984b. “Graphical Perception: Theory, Experimentation, and Application to the Development of Graphical Methods.” *Journal of the American Statistical Association* 79 (387): 531–54. <https://doi.org/10.1080/01621459.1984.10478080>.
- . 1984a. “Graphical Perception: Theory, Experimentation, and Application to the Development of Graphical Methods.” *Journal of the American Statistical Association* 79 (387): 531–54. <https://doi.org/10.1080/01621459.1984.10478080>.
- Croxtan, Frederick E., and Harold Stein. 1932b. “Graphic Comparisons by Bars, Squares, Circles, and Cubes.” *Journal of the American Statistical Association* 27 (177): 54–60. <https://doi.org/10.1080/01621459.1932.10503227>.

- . 1932a. “Graphic Comparisons by Bars, Squares, Circles, and Cubes.” *Journal of the American Statistical Association* 27 (177): 54–60. <https://doi.org/10.1080/01621459.1932.10503227>.
- Croxtan, Frederick E., and Roy E. Stryker. 1927b. “Bar Charts Versus Circle Diagrams.” *Journal of the American Statistical Association* 22 (160): 473–82. <https://doi.org/10.2307/2276829>.
- . 1927a. “Bar Charts Versus Circle Diagrams.” *Journal of the American Statistical Association* 22 (160): 473–82. <https://doi.org/10.2307/2276829>.
- Day, Ross H, and Erica J Stecher. 1991. “Sine of an Illusion.” *Perception* 20 (1): 49–55. <https://doi.org/10.1068/p200049>.
- Dodziuk, Helena. 2016. “Applications of 3D Printing in Healthcare.” *Polish Journal of Cardio-Thoracic Surgery* 3: 283–93. <https://doi.org/10.5114/kitp.2016.62625>.
- Eells, Walter Crosby. 1926a. “The Relative Merits of Circles and Bars for Representing Component Parts.” *Journal of the American Statistical Association* 21 (154): 119–32. <https://doi.org/10.1080/01621459.1926.10502165>.
- . 1926b. “The Relative Merits of Circles and Bars for Representing Component Parts.” *Journal of the American Statistical Association* 21 (154): 119–32. <https://doi.org/10.1080/01621459.1926.10502165>.
- Feinerer, Ingo, and Kurt Hornik. 2024. “Tm: Text Mining Package.” <https://CRAN.R-project.org/package=tm>.
- Fischer, Martin H. 2000. “Do Irrelevant Depth Cues Affect the Comprehension of Bar Graphs?” *Applied Cognitive Psychology* 14 (2): 151–62. [https://doi.org/10.1002/\(SICI\)1099-0720\(200003/04\)14:2%3C151::AID-ACP629%3E3.0.CO;2-Z](https://doi.org/10.1002/(SICI)1099-0720(200003/04)14:2%3C151::AID-ACP629%3E3.0.CO;2-Z).
- Fischer, Martin H., Nele Dewulf, and Robin L. Hill. 2005. “Designing Bar Graphs: Orien-

- tation Matters.” *Applied Cognitive Psychology* 19 (7): 953–62. <https://doi.org/10.1002/acp.1105>.
- Fisher, Samuel H., John V. Dempsey, and Robert T. Marousky. 1997b. “Data Visualization: Preference and Use of Two-Dimensional and Three-Dimensional Graphs.” *Social Science Computer Review* 15 (3): 256–63. <https://doi.org/10.1177/089443939701500303>.
- . 1997a. “Data Visualization: Preference and Use of Two-Dimensional and Three-Dimensional Graphs.” *Social Science Computer Review* 15 (3): 256–63. <https://doi.org/10.1177/089443939701500303>.
- Fleet, Chancey. 2023. “How Tactile Graphics Can Help End Image Poverty.” *MIT Technology Review*, June. <https://www.technologyreview.com/2023/06/15/1074036/ending-image-poverty/>.
- Friendly, Michael, and Howard Wainer. 2021. *A history of data visualization and graphic communication*. Cambridge, Massachusetts: Harvard University Press.
- Graphics, WSJ com News. n.d. “How Does 3-D Printing Work? - WSJ.com.” *The Wall Street Journal*. <http://graphics.wsj.com/3Dprinting>. Accessed September 25, 2025.
- Grinstein, Georges, and Marjan Trutschl. 2001. “High-Dimensional Visualizations.” In: *7th Data Mining Conference-KDD 2001: San Francisco, California*, January.
- Hagerty, Michael, and Michael H. Birnbaum. 1978. “Nonmetric Tests of Ratio Vs. Subtractive Theories of Stimulus Comparison.” *Perception & Psychophysics* 24 (2): 121–29. <https://doi.org/10.3758/BF03199538>.
- Hastie, Trevor J. 2017. “Generalized Additive Models.” In *Statistical Models in s*, 249–307. Routledge.
- Heer, Jeffrey, and Michael Bostock. 2010a. “CHI ’10: CHI Conference on Human Factors in Computing Systems.” In, 203–12. Atlanta Georgia USA: ACM. <https://doi.org/10.1145/1753326.1753391>.

[1145/1753326.1753357](https://doi.org/10.1145/1753326.1753357).

- . 2010b. “Crowdsourcing Graphical Perception: Using Mechanical Turk to Assess Visualization Design.” In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, 203–12. Atlanta Georgia USA: ACM. <https://doi.org/10.1145/1753326.1753357>.
- Henry, Lionel, and Hadley Wickham. 2025. “Rlang: Functions for Base Types and Core r and ‘Tidyverse’ Features.” <https://rlang.r-lib.org>.
- Herzog, A. Regula, and Jerald G. Bachman. 1981. “Effects of Questionnaire Length on Response Quality.” *The Public Opinion Quarterly* 45 (4): 549–59. <https://www.jstor.org/stable/2748903>.
- Hofmann, Heike, Lendie Follett, Mahbubul Majumder, and Dianne Cook. 2012. “Graphical Tests for Power Comparison of Competing Designs.” *IEEE Transactions on Visualization and Computer Graphics* 18 (12): 2441–48. <https://doi.org/10.1109/TVCG.2012.230>.
- Hofmann, Heike, and Marie Vendettuoli. 2013. “Common Angle Plots as Perception-True Visualizations of Categorical Associations.” *IEEE Transactions on Visualization & Computer Graphics*, no. 12: 2297–2305. <https://doi.org/10.1109/TVCG.2013.140>.
- Hollands, J. G., and Brian P. Dyre. 2000. “Bias in Proportion Judgments: The Cyclical Power Model.” *Psychological Review* 107 (3): 500–524. <https://doi.org/10.1037/0033-295X.107.3.500>.
- Holloway, Leona, Kim Marriott, Matthew Butler, and Samuel Reinders. 2019. “ASSETS ’19: The 21st International ACM SIGACCESS Conference on Computers and Accessibility.” In, 183–95. Pittsburgh PA USA: ACM. <https://doi.org/10.1145/3308561.3353790>.
- Horst, Allison Marie, Alison Presmanes Hill, and Kristen B Gorman. 2020. *Palmerpenguins: Palmer Archipelago (Antarctica) Penguin Data*. <https://doi.org/10.5281/zenodo>.

3960218.

- Huff, Darrell. 1954. *How to Lie with Statistics*. W. W. Norton and Company.
- Hughes, Brian M. n.d. “Just Noticeable Differences in 2D and 3D Bar Charts: A Psychophysical Analysis of Chart Readability.”
- Hullman, Jessica, Xiaoli Qiao, Michael Correll, Alex Kale, and Matthew Kay. 2019. “In Pursuit of Error: A Survey of Uncertainty Visualization Evaluation.” *IEEE Transactions on Visualization and Computer Graphics* 25 (1): 903–13. <https://doi.org/10.1109/TVCG.2018.2864889>.
- Huron, Samuel, Till Nagel, Lora Oehlberg, and Wesley Willett, eds. 2023. *Making with Data: Physical Design and Craft in a Data-Driven World*. First edition. AK Peters Visualization Series. Boca Raton: AK Peters : CRC Press.
- Jeong, Jongwook, Ayoung Choi, Dongwoo Kang, and Youn Kyu Lee. 2025. “A Comparative Study of 2D Vs. 3D Chart Visualizations in Virtual Reality.” *Journal of Visualization* 28 (1): 239–53. <https://doi.org/10.1007/s12650-024-01033-6>.
- Katsioloudis, Petros, and Mildred Jones. 2018. “A Comparative Analysis of Holographic, 3D-Printed, and Computer-Generated Models: Implications for Engineering Technology Students’ Spatial Visualization Ability.” *Journal of Technology Education* 29 (2): 36–53. <https://doi.org/10.21061/jte.v29i2.a.3>.
- Kingdom, Frederick A. A., and Nicolaas Prins. 2016. *Psychophysics: A Practical Introduction*. Second edition. Amsterdam: Elsevier/Academic Press.
- Kintel, Marius. 2023. “OpenSCAD. OpenSCAD.org.” July 17, 2023. <http://openscad.org/documentation.html>.
- Kraus, Matthias, Katrin Angerbauer, Juri Buchmüller, Daniel Schweitzer, Daniel A. Keim, Michael Sedlmair, and Johannes Fuchs. 2020. “CHI ’20: CHI Conference on Human

- Factors in Computing Systems.” In, 1–14. Honolulu HI USA: ACM. <https://doi.org/10.1145/3313831.3376675>.
- Levy, Ellen, Jeff Zacks, Barbara Tversky, and Diane Schiano. 1996. “The SIGCHI Conference.” In, 42–49. Vancouver, British Columbia, Canada: ACM Press. <https://doi.org/10.1145/238386.238400>.
- Liu, Yang, and Jeffrey Heer. 2018a. “Somewhere Over the Rainbow: An Empirical Assessment of Quantitative Colormaps.” In *Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems*, 1–12. Montreal QC Canada: ACM. <https://doi.org/10.1145/3173574.3174172>.
- . 2018b. “CHI ’18: CHI Conference on Human Factors in Computing Systems.” In, 1–12. Montreal QC Canada: ACM. <https://doi.org/10.1145/3173574.3174172>.
- Loy, Adam. 2021. “Bringing Visual Inference to the Classroom.” *Journal of Statistics and Data Science Education* 29 (2): 171–82. <https://doi.org/10.1080/26939169.2021.1920866>.
- Maineri, Angelica M., Ivano Bison, and Ruud Luijkx. 2021. “Slider Bars in Multi-Device Web Surveys.” *Social Science Computer Review* 39 (4): 573–91. <https://doi.org/10.1177/0894439319879132>.
- Mandal, B. N., Rajender Parsad, and Sukanta Dash. 2020. “Incomplete Split-Plot Designs: Construction and Analysis.” *Statistics & Probability Letters* 166 (November): 108869. <https://doi.org/10.1016/j.spl.2020.108869>.
- Margaret, Mackisack. 1994. “What Is the Use of Experiments Conducted By Statistics Students?” *Journal of Statistics Education* 2 (1): 2. <https://doi.org/10.1080/10691898.1994.11910461>.
- McGowan, Herle M. 2011. “Planning a Comparative Experiment in Educational Set-

- tings.” *Journal of Statistics Education* 19 (2): 4. <https://doi.org/10.1080/10691898.2011.11889612>.
- Melody Carswell, C., Sylvia Frankenberger, and Donald Bernhard. 1991. “Graphing in Depth: Perspectives on the Use of Three-Dimensional Graphs to Represent Lower-Dimensional Data.” *Behaviour & Information Technology* 10 (6): 459–74. <https://doi.org/10.1080/01449299108924304>.
- Microsoft Corporation. 2018. “Microsoft Excel.”
- Molina Lopez, Elena, Carolina Middel Soria, and Pere-Pau Vázquez. 2023. “Effect of Color Palettes in Heatmaps Perception: A Study,” 5 pages. <https://doi.org/10.2312/EVS.20231039>.
- Morgan-Wall, Tyler. 2019. “Introducing 3D Ggplots with Rayshader.” Blog. *Tyler Morgan-Wall*. https://www.tylermw.com/posts/data_visualization/3d-ggplots-with-rayshader.html.
- . 2024. “Rayshader: Create Maps and Visualize Data in 2D and 3D.” <https://CRAN.R-project.org/package=rayshader>.
- Mozilla Foundation. 2023. “WebGL: 2D and 3D Graphics for the Web.” July 12, 2023. https://developer.mozilla.org/en-US/docs/Web/API/WebGL_API.
- Muff, Julian Louis, Tobias Heye, Florian Markus Thieringer, and Philipp Brantner. 2022. “Clinical Acceptance of Advanced Visualization Methods: A Comparison Study of 3D-Print, Virtual Reality Glasses, and 3D-Display.” *3D Printing in Medicine* 8 (January): 5. <https://doi.org/10.1186/s41205-022-00133-z>.
- Munzner, Tamara. n.d. “Visualization Design and Analysis: Abstractions, Principles, and Methods.”
- Murdoch, Duncan, and Daniel Adler. 2023. *Rgl: 3D Visualization Using OpenGL*.

- Peña, Alyssa, Eric Ragan, and Lane Harrison. 2020. “Memorability of Enhanced Informational Graphics.” *Interdisciplinary Journal of Signage and Wayfinding* 4 (1). <https://doi.org/10.15763/issn.2470-9670.2020.v4.i1.a54>.
- R Core Team. 2024. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org/>.
- Reda, Khairi, and Michael E. Papka. 2019. “2019 IEEE Visualization Conference (VIS).” In, 271–75. <https://doi.org/10.1109/VISUAL.2019.8933760>.
- Rice, Kiegan, Heike Hofmann, Nola Du Toit, and Edward Mulrow. 2024. “Testing Perceptual Accuracy in a U.S. General Population Survey Using Stacked Bar Charts.” *Journal of Data Science*, 280–97. <https://doi.org/10.6339/24-JDS1121>.
- Ruud, Paul A., Daniel Schunk, and Joachim K. Winter. 2014. “Uncertainty Causes Rounding: An Experimental Study.” *Experimental Economics* 17 (3): 391–413. <https://doi.org/10.1007/s10683-013-9374-8>.
- Sarkar, Deepayan. 2008. “Lattice: Multivariate Data Visualization with r.” <http://lmdvr.r-forge.r-project.org>.
- “SAS 9.4 ODS Graphics: Procedures Guide, Sixth Edition.” n.d.
- Satyanarayan, Arvind, Dominik Moritz, Kanit Wongsuphasawat, and Jeffrey Heer. 2017. “Vega-Lite: A Grammar of Interactive Graphics.” *IEEE Transactions on Visualization and Computer Graphics* 23 (1): 341–50. <https://doi.org/10.1109/TVCG.2016.2599030>.
- Schmid, Calvin F. 1983. *Statistical graphics: design principles and practices*. A Wiley-interscience publication. New York: Wiley.
- Siegrist, Michael. 1996. “The Use or Misuse of Three-Dimensional Graphs to Represent Lower-Dimensional Data.” *Behaviour & Information Technology* 15 (2): 96–100. <https://doi.org/10.1080/014492996120300>.

- Sievert, Carson. 2020. “Interactive Web-Based Data Visualization with r, Plotly, and Shiny.” <https://plotly-r.com>.
- Spence, Ian. 1990. “Visual Psychophysics of Simple Graphical Elements.” *Journal of Experimental Psychology: Human Perception and Performance* 16: 683–92. <https://doi.org/10.1037/0096-1523.16.4.683>.
- Stevens, S. S., and Geraldine Stevens. 1986. *Psychophysics: Introduction to Its Perceptual, Neural, and Social Prospects*. New Brunswick, U.S.A: Transaction Books.
- Stewart, Brandie M., Jessica M. Cipolla, and Lisa A. Best. 2009. “Extraneous Information and Graph Comprehension: Implications for Effective Design Choices.” Edited by M. Henderson. *Campus-Wide Information Systems* 26 (3): 191–200. <https://doi.org/10.1108/10650740910967375>.
- Stroup, Walter W, George A Milliken, Elizabeth A Claassen, and Russell D Wolfinger. n.d. *SAS for Mixed Models: Introduction and Basic Applications*.
- Tarr, Michael J, and David J Kriegman. 2001. “What Defines a View?” *Vision Research* 41 (15): 1981–2004. [https://doi.org/10.1016/S0042-6989\(01\)00024-4](https://doi.org/10.1016/S0042-6989(01)00024-4).
- Thingiverse.com. n.d. “Chart: Sales of 3-D Printers Under \$5,000 by WSJ.” *Thingiverse*. <https://www.thingiverse.com/thing:365994>. Accessed September 23, 2025.
- Tintle, Nathan, Beth L Chance, George W Cobb, Allan J Rossman, Soma Roy, Todd Swanson, and Jill VanderStoep. 2021. *Introduction to Statistical Investigations*.
- Tractinsky, Noam, and Joachim Meyer. 1999. “Chartjunk or Goldgraph? Effects of Presentation Objectives and Content Desirability on Information Presentation.” *MIS Quarterly* 23 (3): 397–420. <https://doi.org/10.2307/249469>.
- Tufte, Edward R. 2001. *The Visual Display of Quantitative Information*. 2nd ed. Cheshire, Conn: Graphics Press.

- Tukey, John W. 1965. “The Technical Tools of Statistics.” *The American Statistician* 19 (2): 23–28. <https://doi.org/10.2307/2682374>.
- V, Sundaram, Vidhya Lakshmi Sivakumar, Anand Raju, Saravanan A, and Titus Richard. 2023. “The Role of 3D Printing in Engineering and Its Historical Perspective in Civil Engineering - A Review.” *International Journal of Civil Engineering* 10 (12): 58–68. <https://doi.org/10.14445/23488352/IJCE-V10I12P107>.
- Vanderplas, Susan, Dianne Cook, and Heike Hofmann. 2020a. “Testing Statistical Charts: What Makes a Good Graph?” *Annual Review of Statistics and Its Application* 7 (1): annurev-statistics-031219-041252. <https://doi.org/ggnv7t>.
- . 2020b. “Testing Statistical Charts: What Makes a Good Graph?” *Annual Review of Statistics and Its Application* 7 (1): 61–88. <https://doi.org/10.1146/annurev-statistics-031219-041252>.
- VanderPlas, Susan, and Heike Hofmann. 2015a. “Signs of the Sine Illusion—Why We Need to Care.” *Journal of Computational and Graphical Statistics* 24 (4): 1170–90. <https://doi.org/10.1080/10618600.2014.951547>.
- . 2015b. “Signs of the Sine Illusion—Why We Need to Care.” *Journal of Computational and Graphical Statistics* 24 (4): 1170–90. <https://doi.org/10.1080/10618600.2014.951547>.
- VanderPlas, Susan, Goluch C Ryan, and Heike Hofmann. 2019. “Framed! Reproducing and Revisiting 150-Year-Old Charts.” *Journal of Computational and Graphical Statistics* 28 (3): 620–34. <https://doi.org/10.1080/10618600.2018.1562937>.
- Veit, Clairice T. n.d. “Ratio and Subtractive Processes in Psychophysical Judgment.”
- von Huhn, R. 1927. “Further Studies in the Graphic Use of Circles and Bars.” *Journal of the American Statistical Association* 22 (157): 31–36. <https://doi.org/10.1080/01621459>.

1927.10502938.

- Wang, Xiyao, Lonni Besançon, Mehdi Ammi, and Tobias Isenberg. 2022. “Understanding Differences Between Combinations of 2D and 3D Input and Output Devices for 3D Data Visualization.” *International Journal of Human-Computer Studies* 163 (July): 102820. <https://doi.org/10.1016/j.ijhcs.2022.102820>.
- Wickham, Hadley. 2010. “A Layered Grammar of Graphics.” *Journal of Computational and Graphical Statistics* 19 (1): 3–28. <https://www.jstor.org/stable/25651297>.
- . 2016. “Ggplot2: Elegant Graphics for Data Analysis.” <https://ggplot2.tidyverse.org>.
- Wiederich, T. 2023. “Evaluating Perceptual Judgements on 3D Printed Bar Charts,” May. https://twiedrw.github.io/3d_graphical_perception/SDSS_presentation/index.html.
- Wiederich, T, and S VanderPlas. n.d. “Placeholder: A Comparative Experiment on Projections of 2D and 3D Bar Charts.”
- Wilke, Claus O. 2020. “Don’t Go 3D.” In *Fundamentals of Data Visualization*. O’Reilly Media. <https://clauswilke.com/dataviz/no-3d.html>.
- Wilkinson, Leland, and Graham Wills. 2005. *The Grammar of Graphics*. 2nd ed. Statistics and Computing. New York: Springer.
- Wood, S. N. 2017. *Generalized Additive Models: An Introduction with r*. 2nd ed. Chapman; Hall/CRC.
- Wood, S. N. 2003. “Thin-Plate Regression Splines.” *Journal of the Royal Statistical Society (B)* 65 (1): 95–114.
- . 2004. “Stable and Efficient Multiple Smoothing Parameter Estimation for Generalized Additive Models.” *Journal of the American Statistical Association* 99 (467): 673–86.

- Zacks, Jeff, Ellen Levy, Barbara Tversky, and Diane J. Schiano. 1998. “Reading Bar Graphs: Effects of Extraneous Depth Cues and Graphical Context.” *Journal of Experimental Psychology: Applied* 4 (2): 119–38. <https://doi.org/10.1037/1076-898X.4.2.119>.
- Zhang, Zhao, Xiwu Shang, Guoping Li, and Guozhong Wang. 2023. “Just Noticeable Difference Model for Images with Color Sensitivity.” *Sensors* 23 (5): 2634. <https://doi.org/10.3390/s23052634>.